

A Journey Into Node.js Internals

Tamar Twena-Stern

Tamar Twena-Stern



- Software Engineer - manager and architect
- Architect @PaloAltoNetworks
- Was a CTO of my own startup
- Passionate about Node.js !
- Twitter: **@SternTwena**

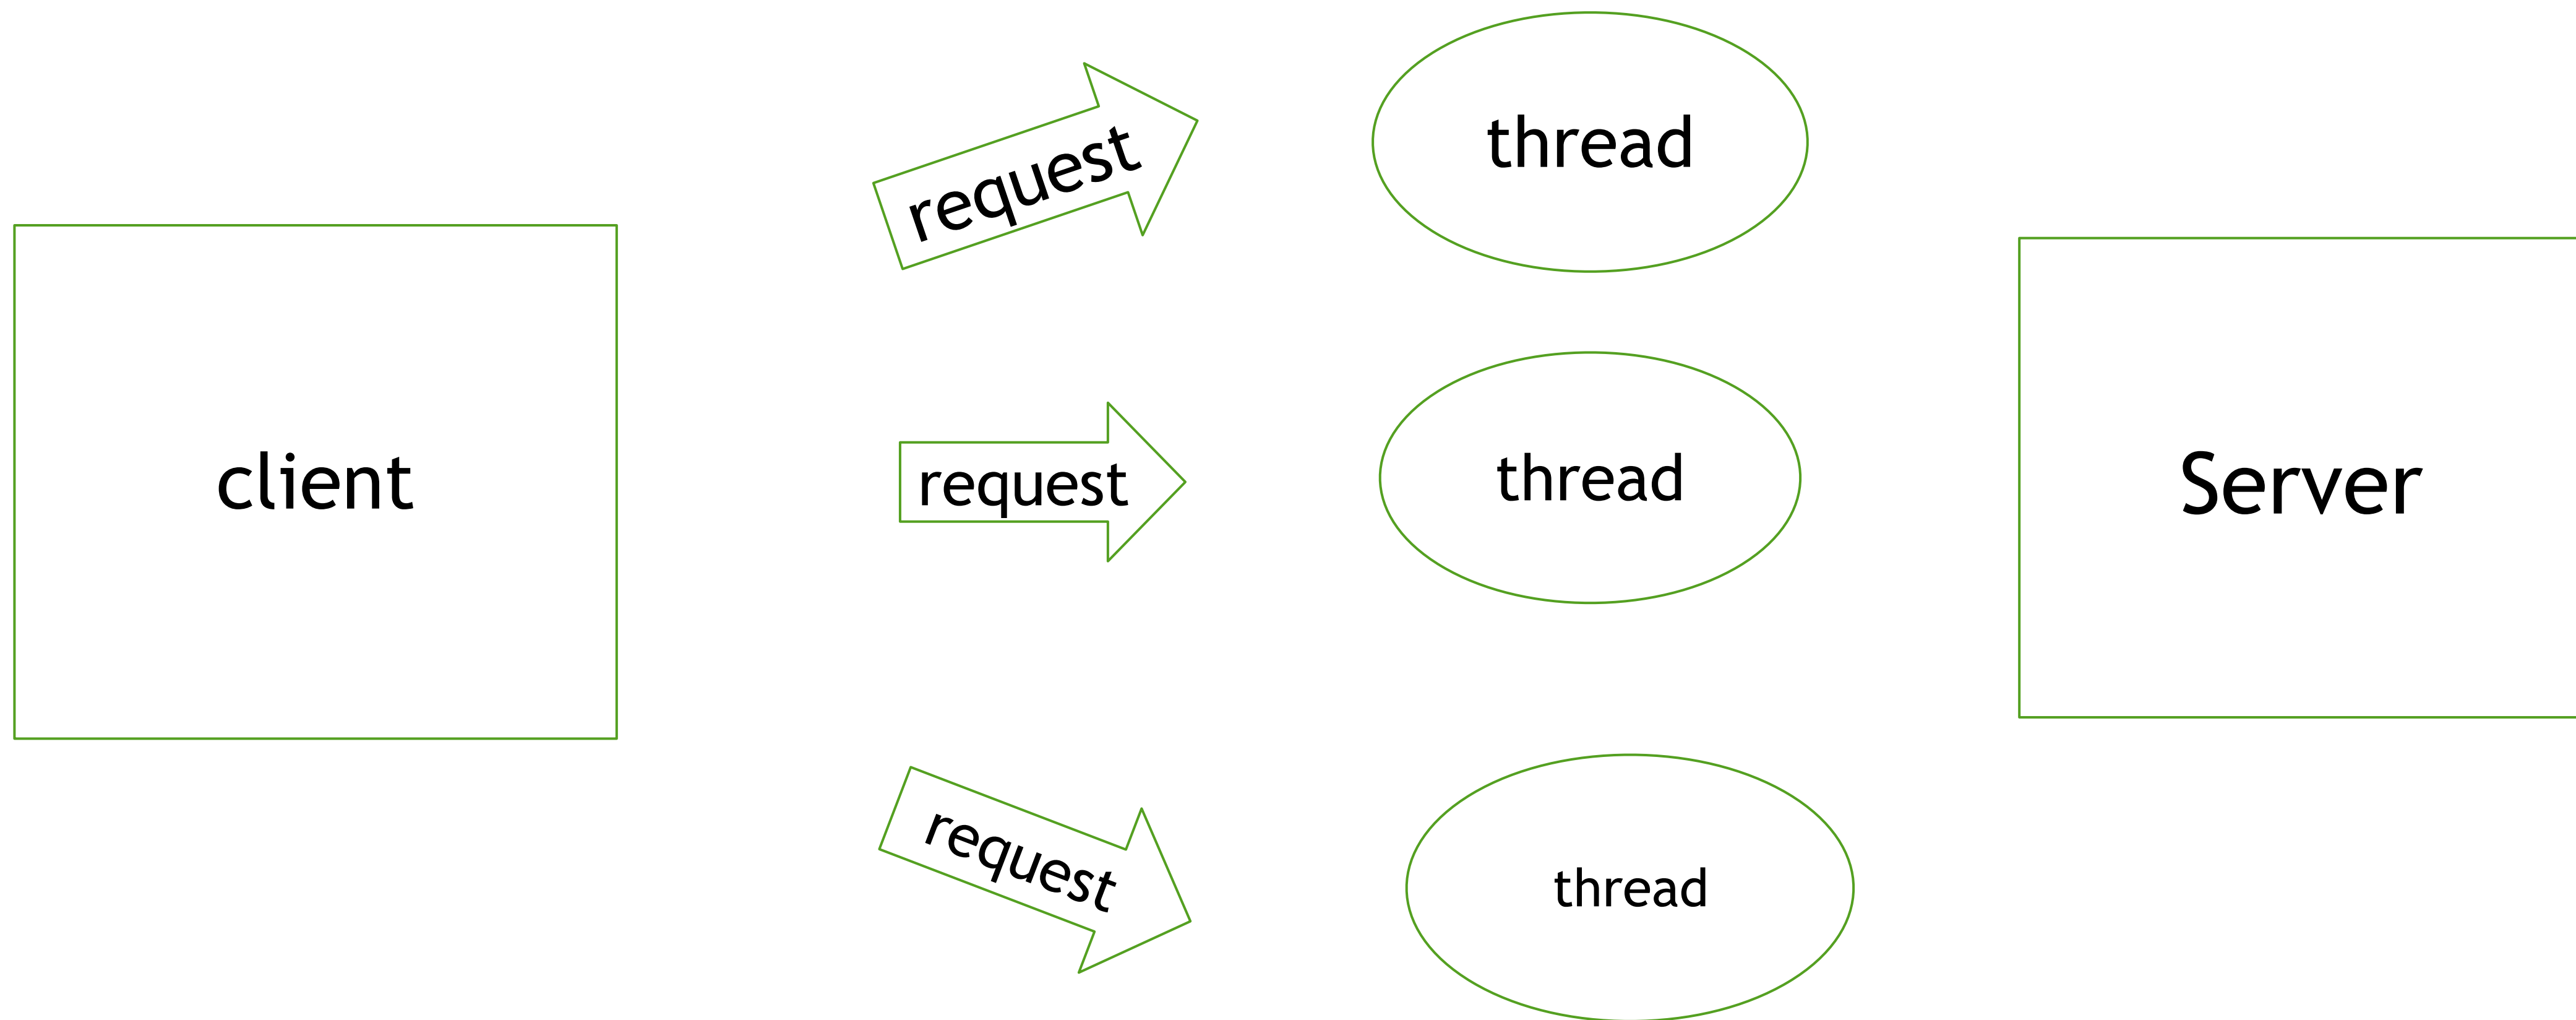
Tamar Tena-Stern

- On Maternity Leave
- Have 3 kids
- Loves to play my violin
- Javascript Israel community leader



Introduction

Traditional Approach

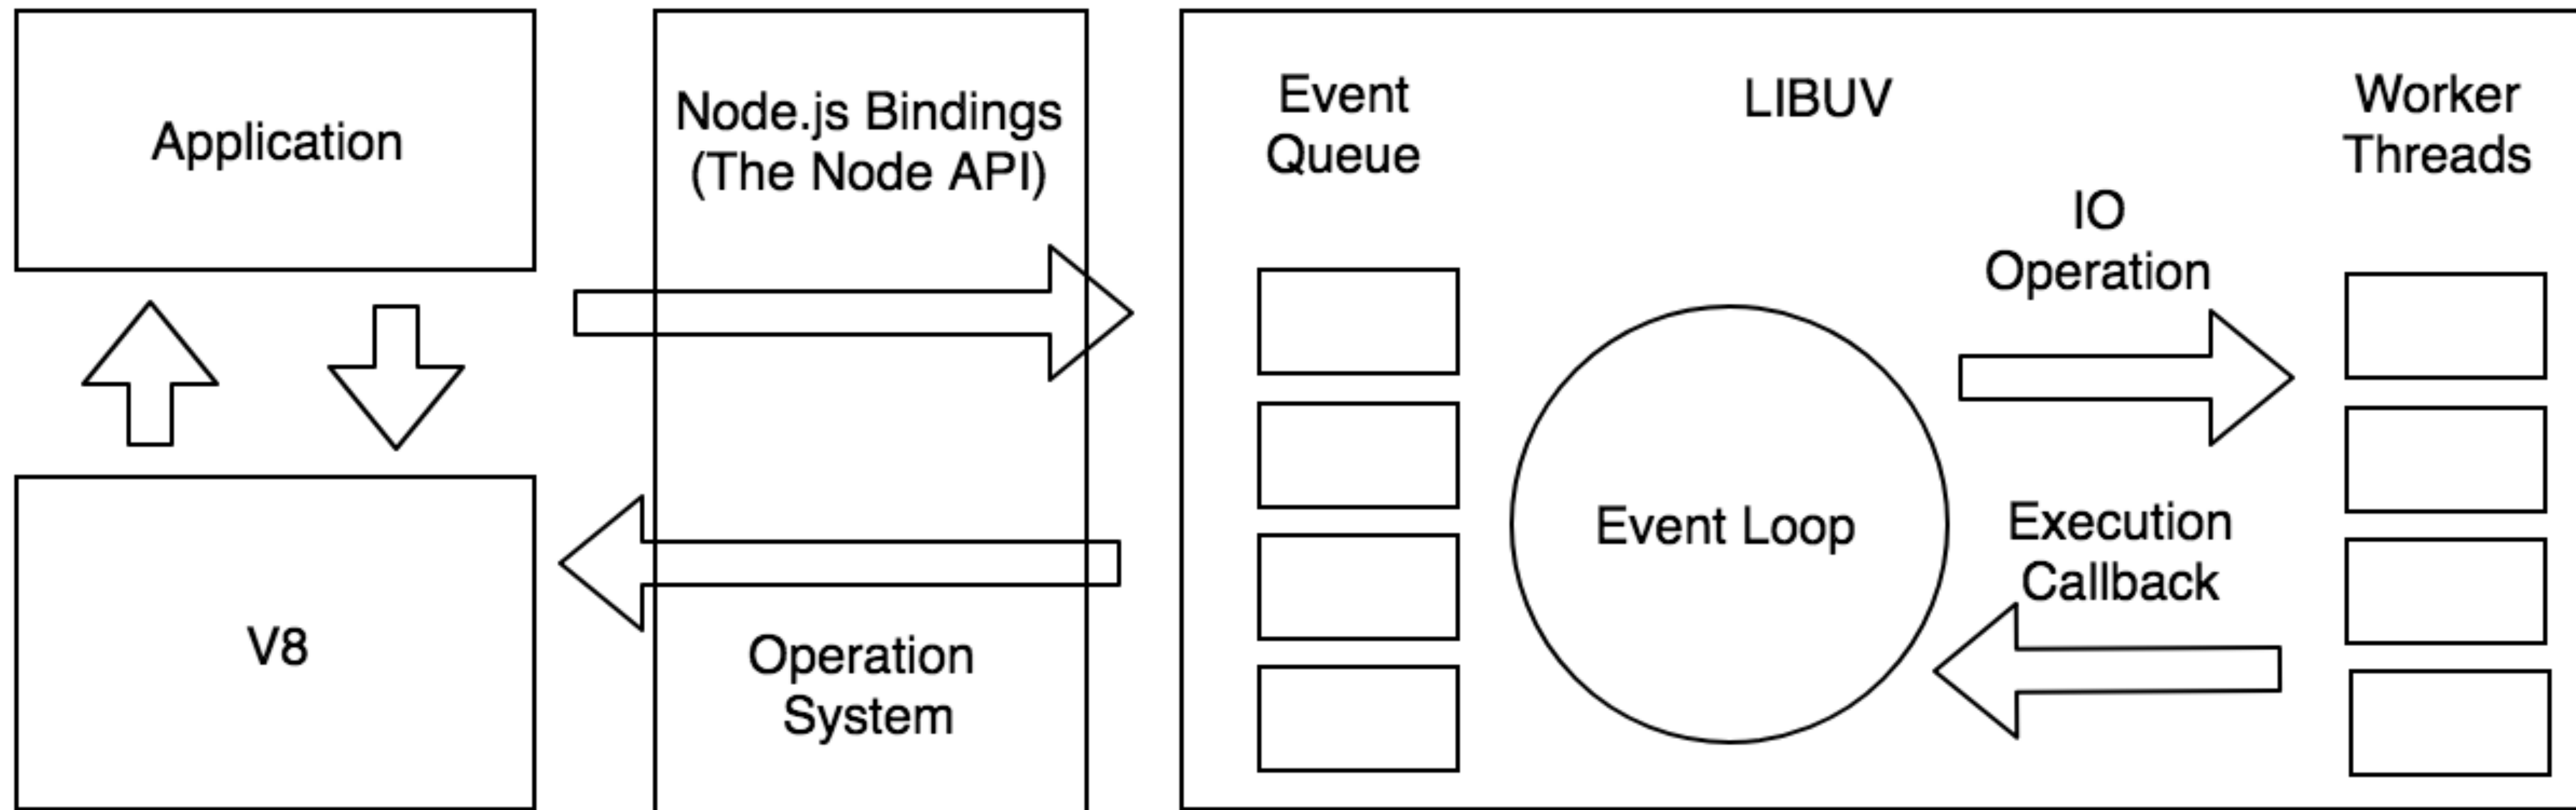


Problems

- The system allocates CPU and memory resources for every new thread
- When the system is stressed - overhead of thread scheduling and context switching
- The system waste resources for allocating threads instead of doing actual work

Node.js Architecture

Node.js Architecture - High Level

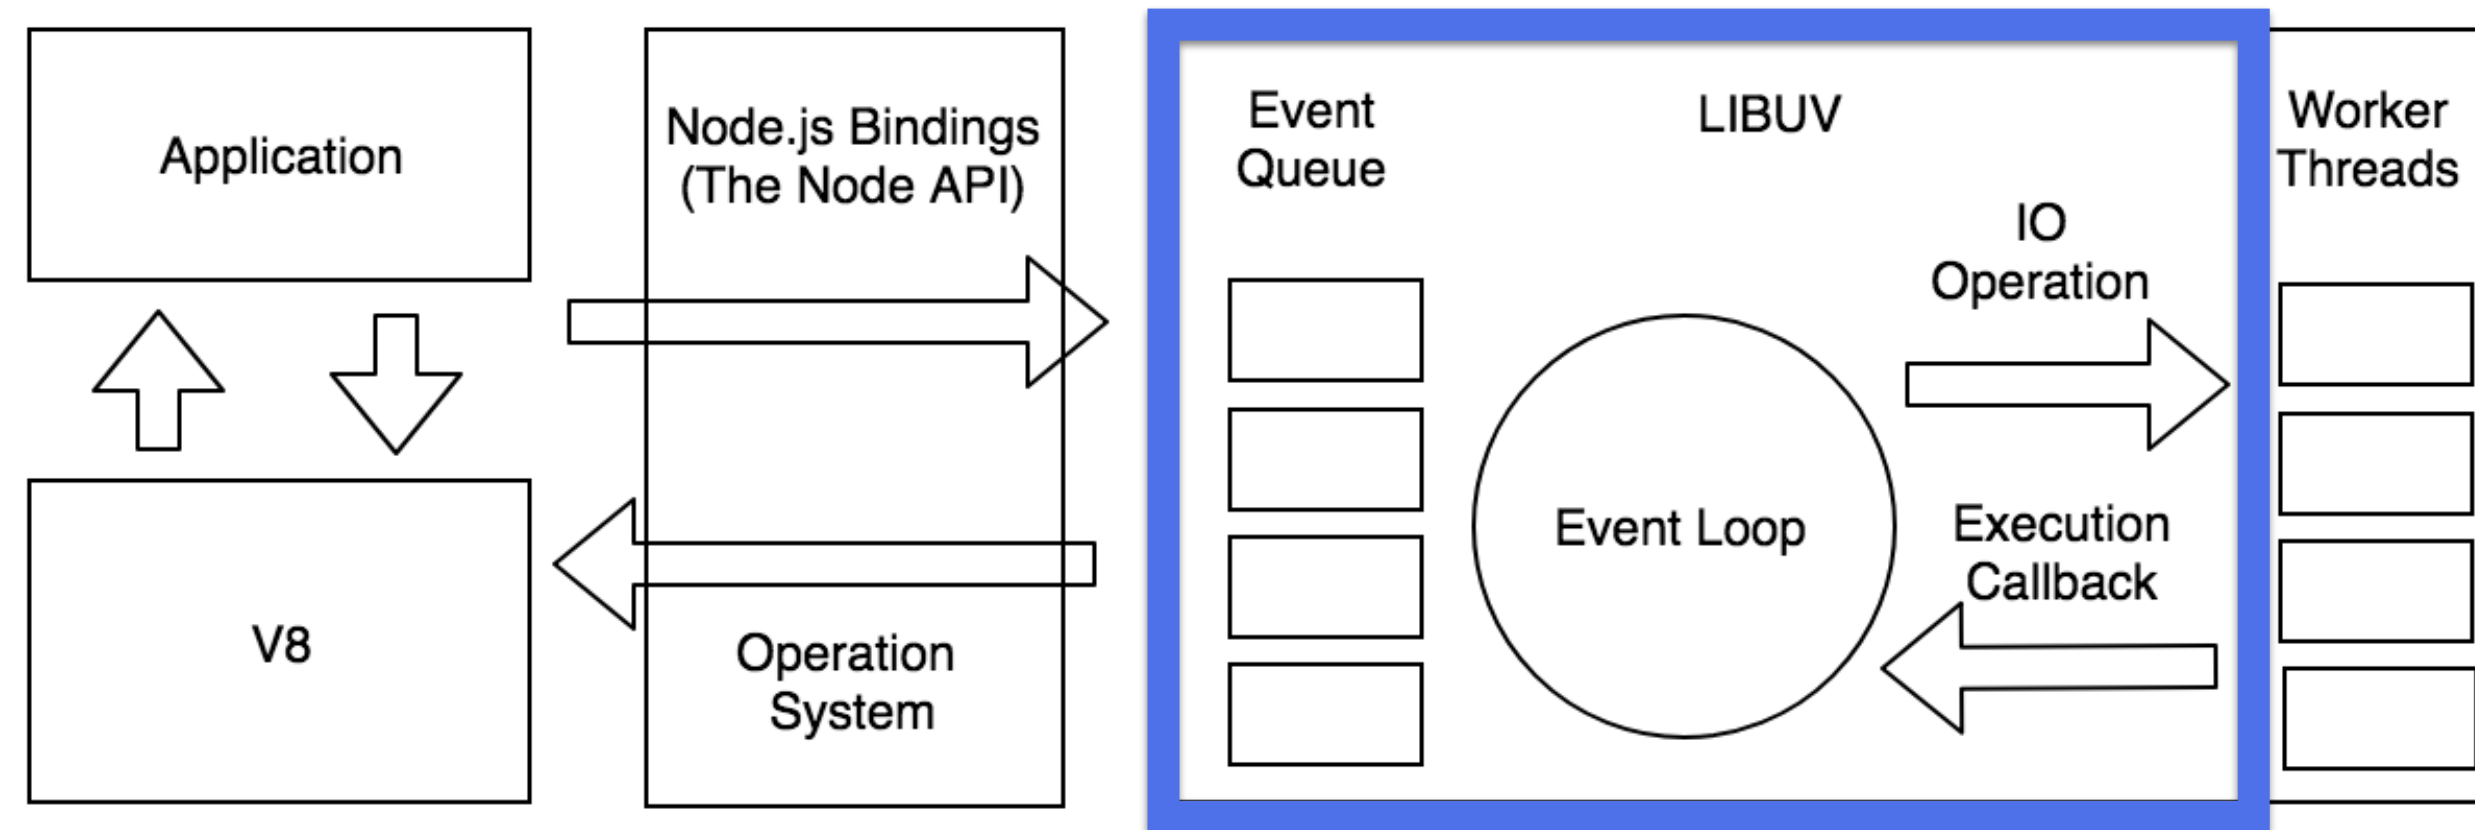


Now, Lets Get Into The Details

Single Threaded ?

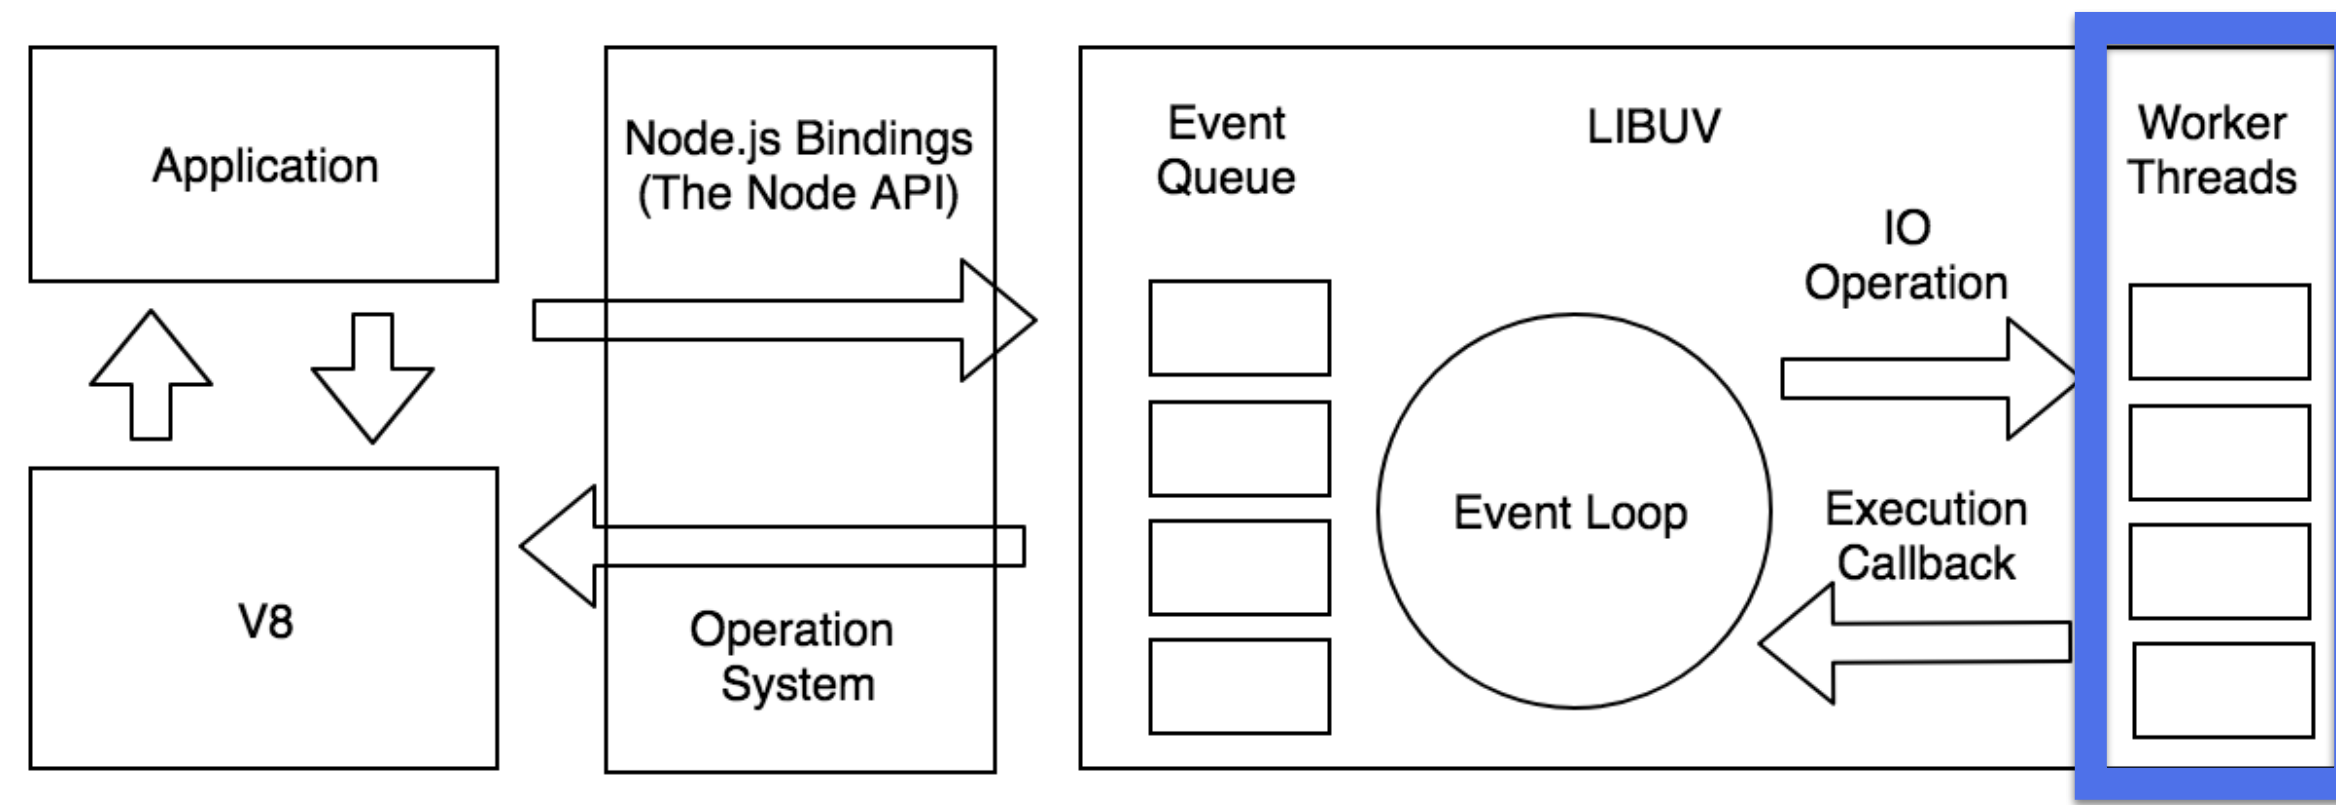
- Not really single threaded
- Several threaded :
 - Event Loop
 - The workers thread pool

Event Loop Thread



- Every request registers a callback which executes immediately
- The event loop execute JavaScript callbacks
- Offloads I/O operations to worker thread pool.
- Handle callbacks for asynchronous I/O operations from multiple requests.

Worker Thread Pool

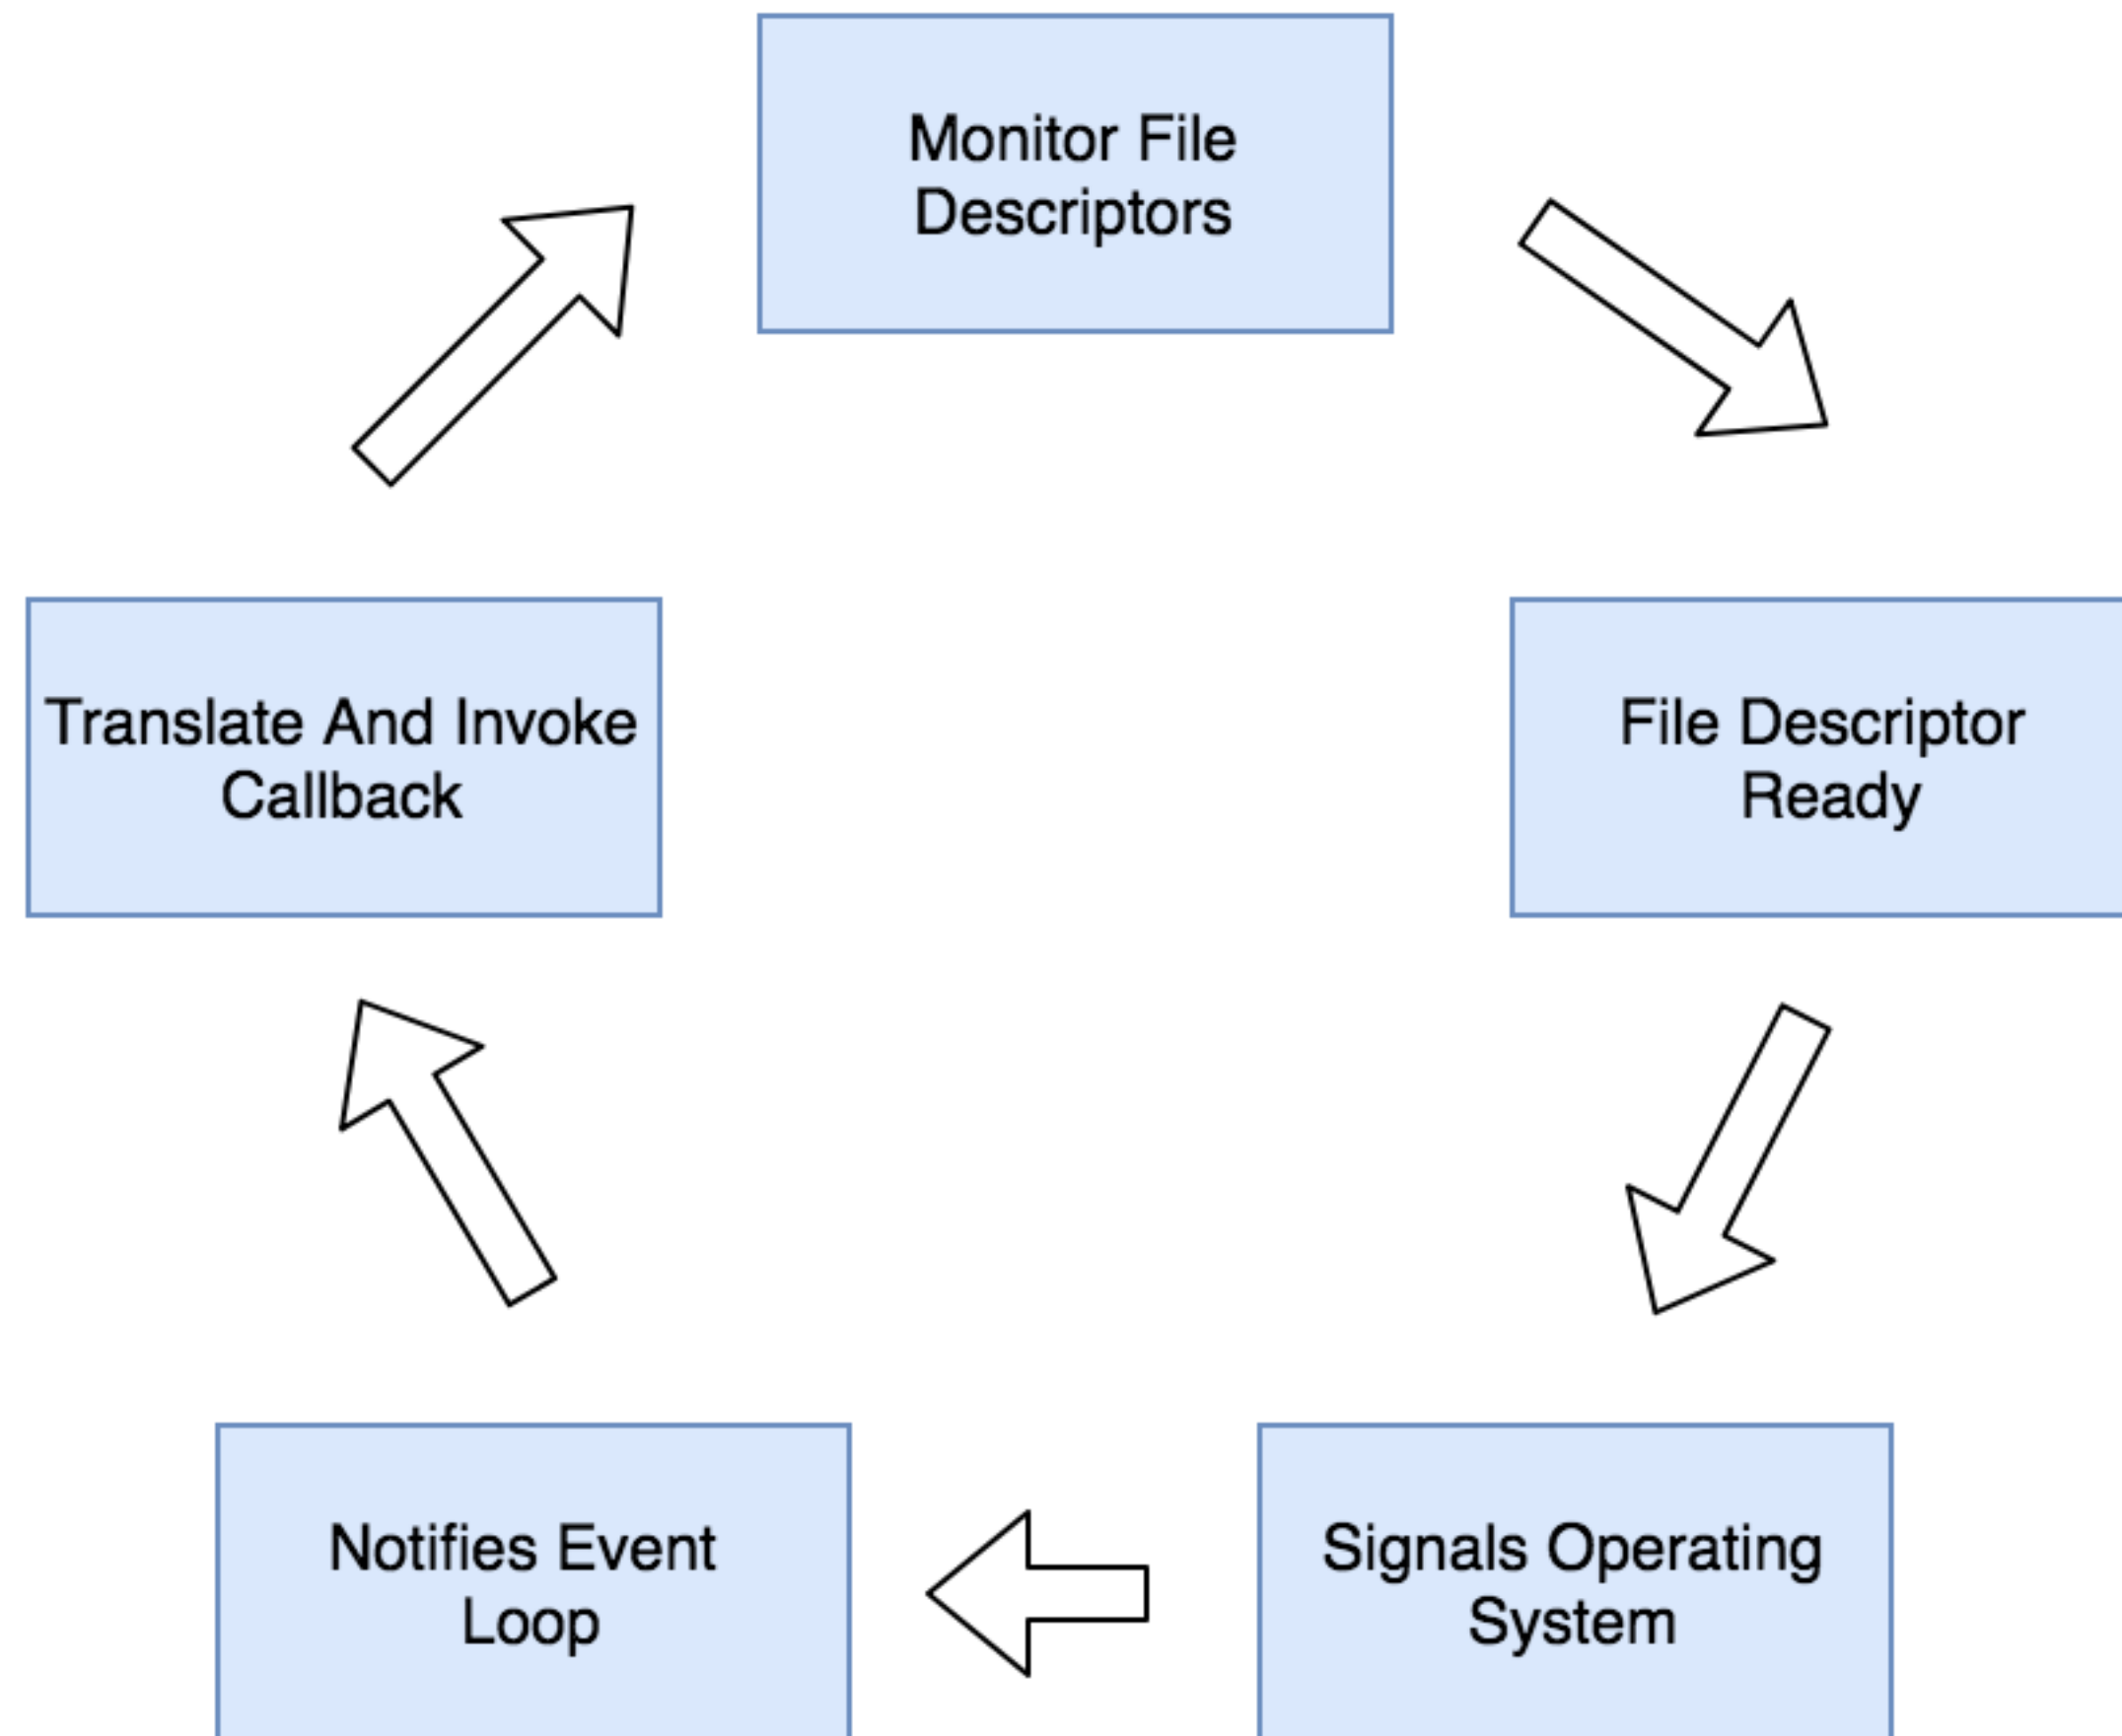


- Thread pool to perform heavy operations
 - I/O
 - CPU intensive operations
- Bounded by fixed capacity
- A node module can submit a task to libUV API

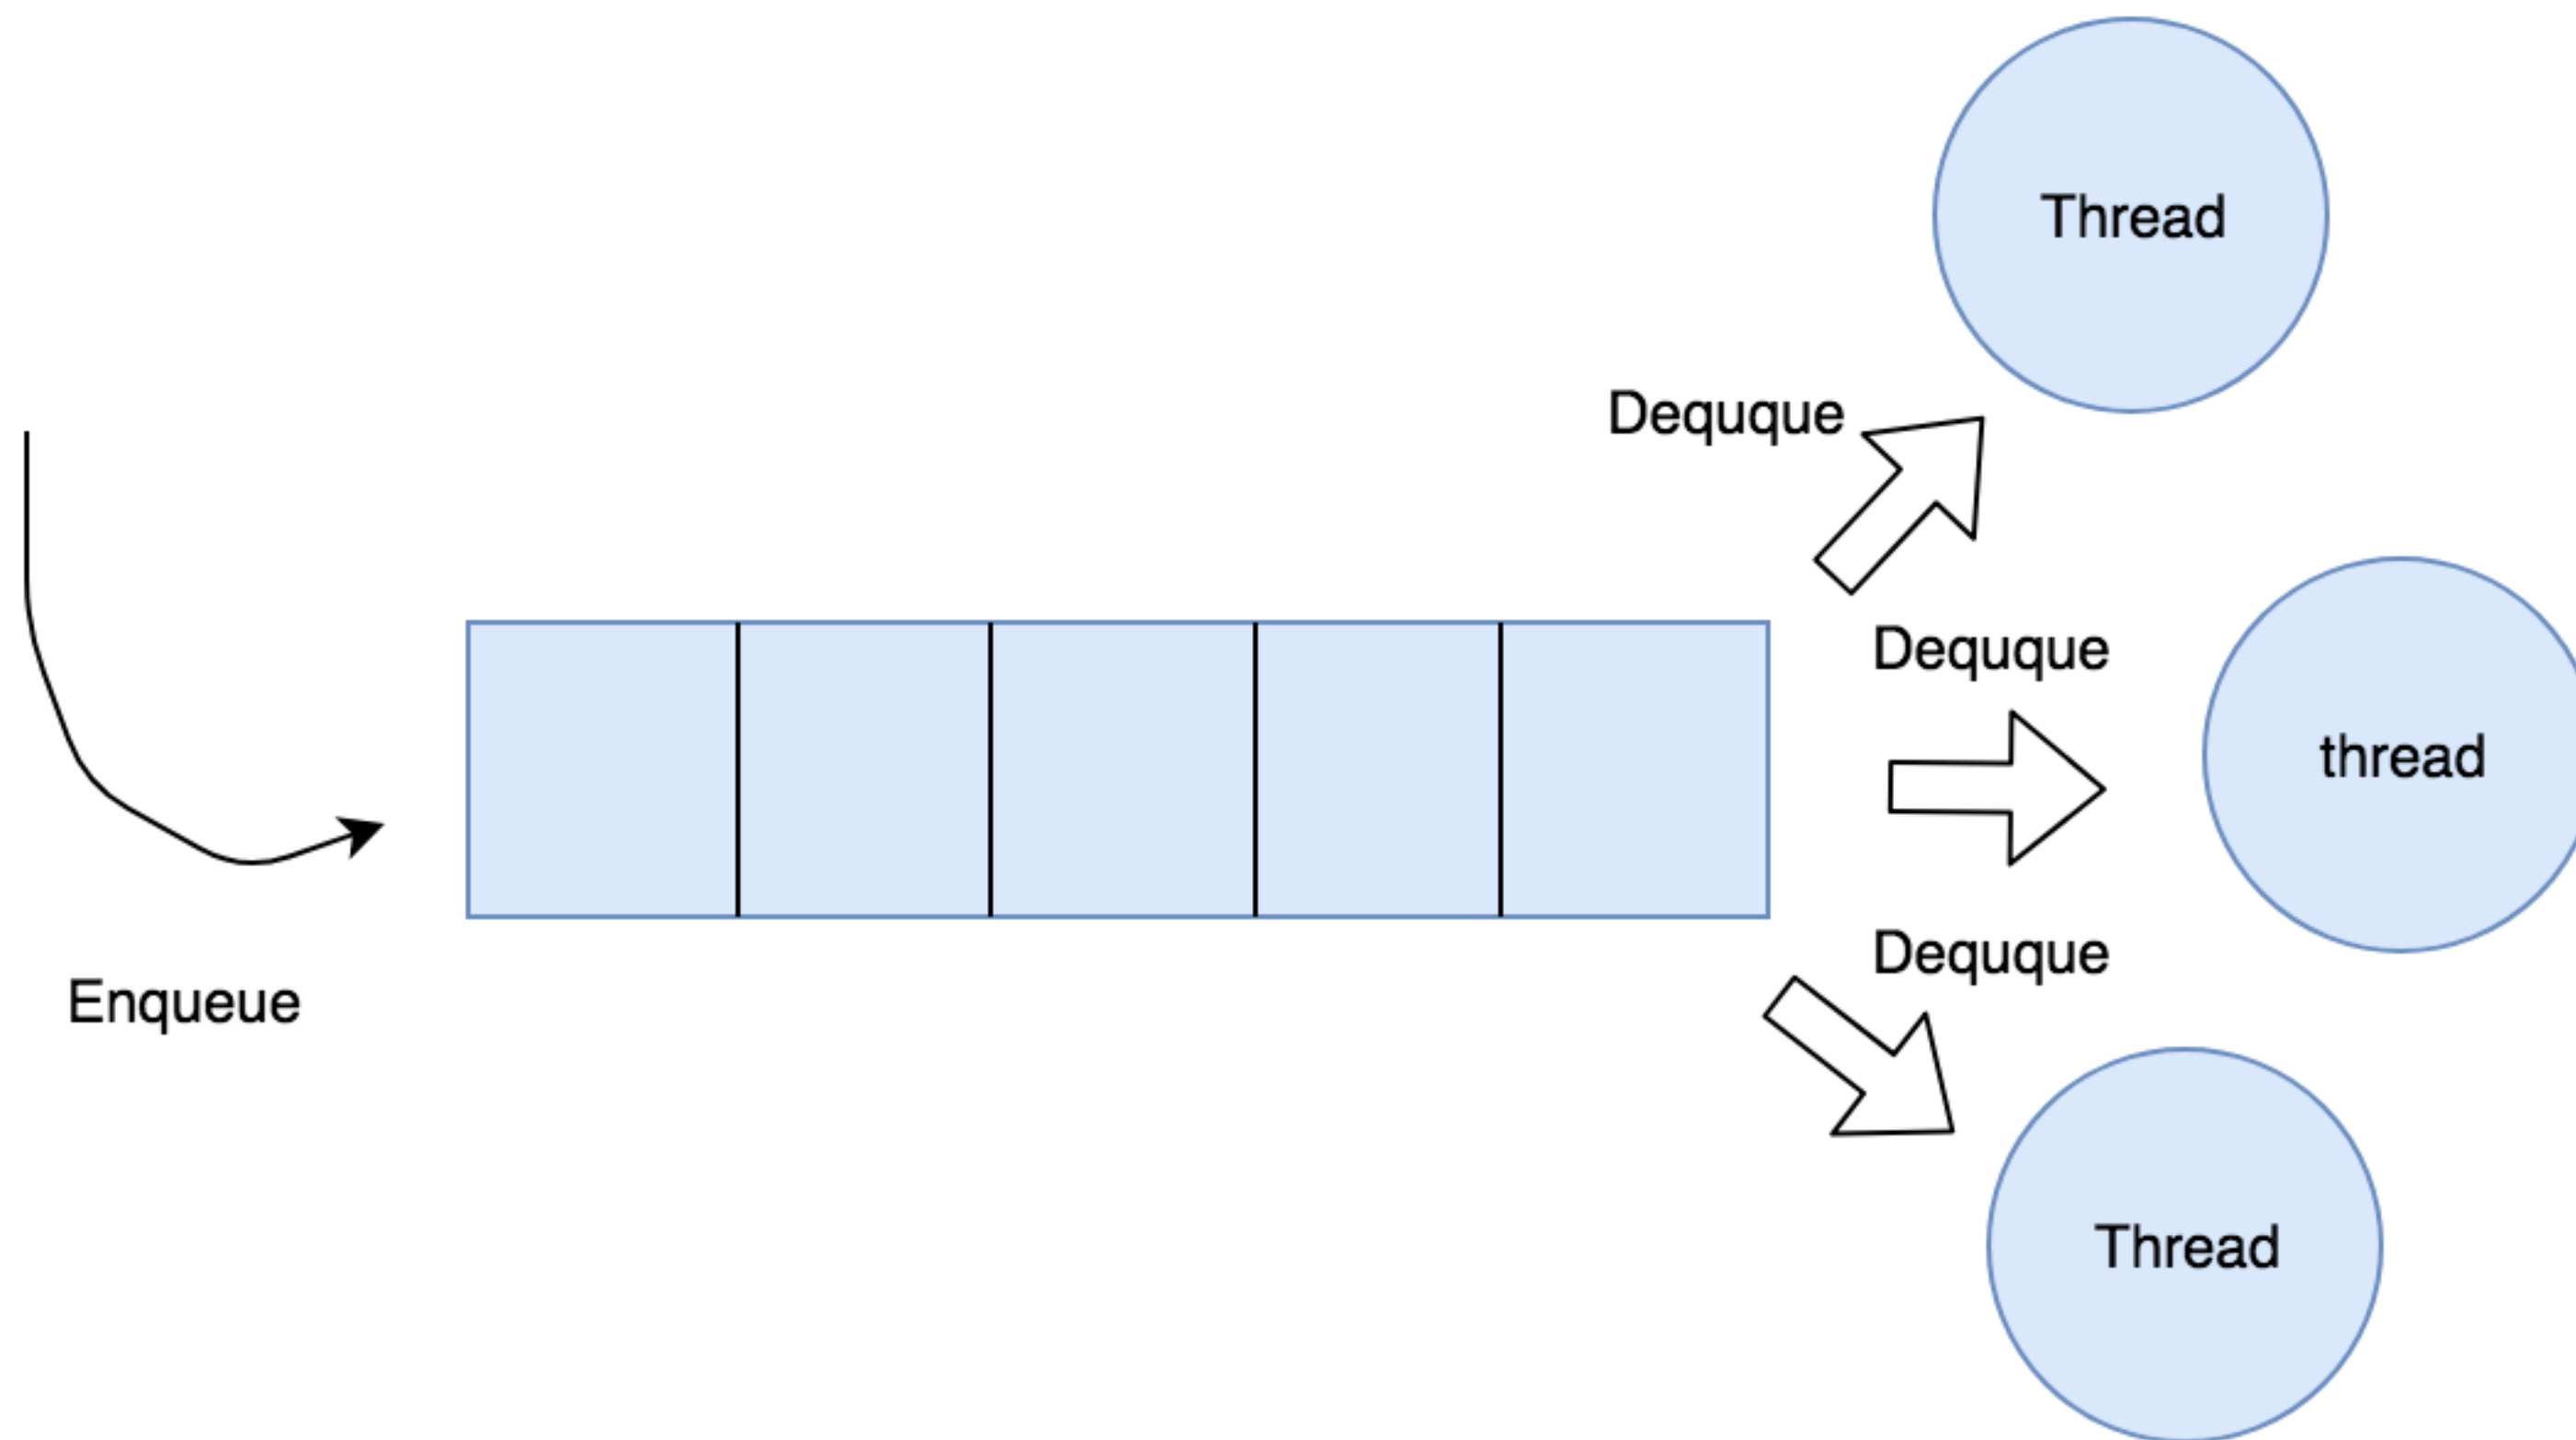
Submitting A Request To The Worker Pool

- Use a set of ‘basic’ modules that work with the event loop
 - Examples:
 - Fs
 - Dns
 - Crypto
 - And more
- Submit a task to libUV using c++ add-on

Event Loop Implemented With A Queue ?



How Worker Pool Implemented ?



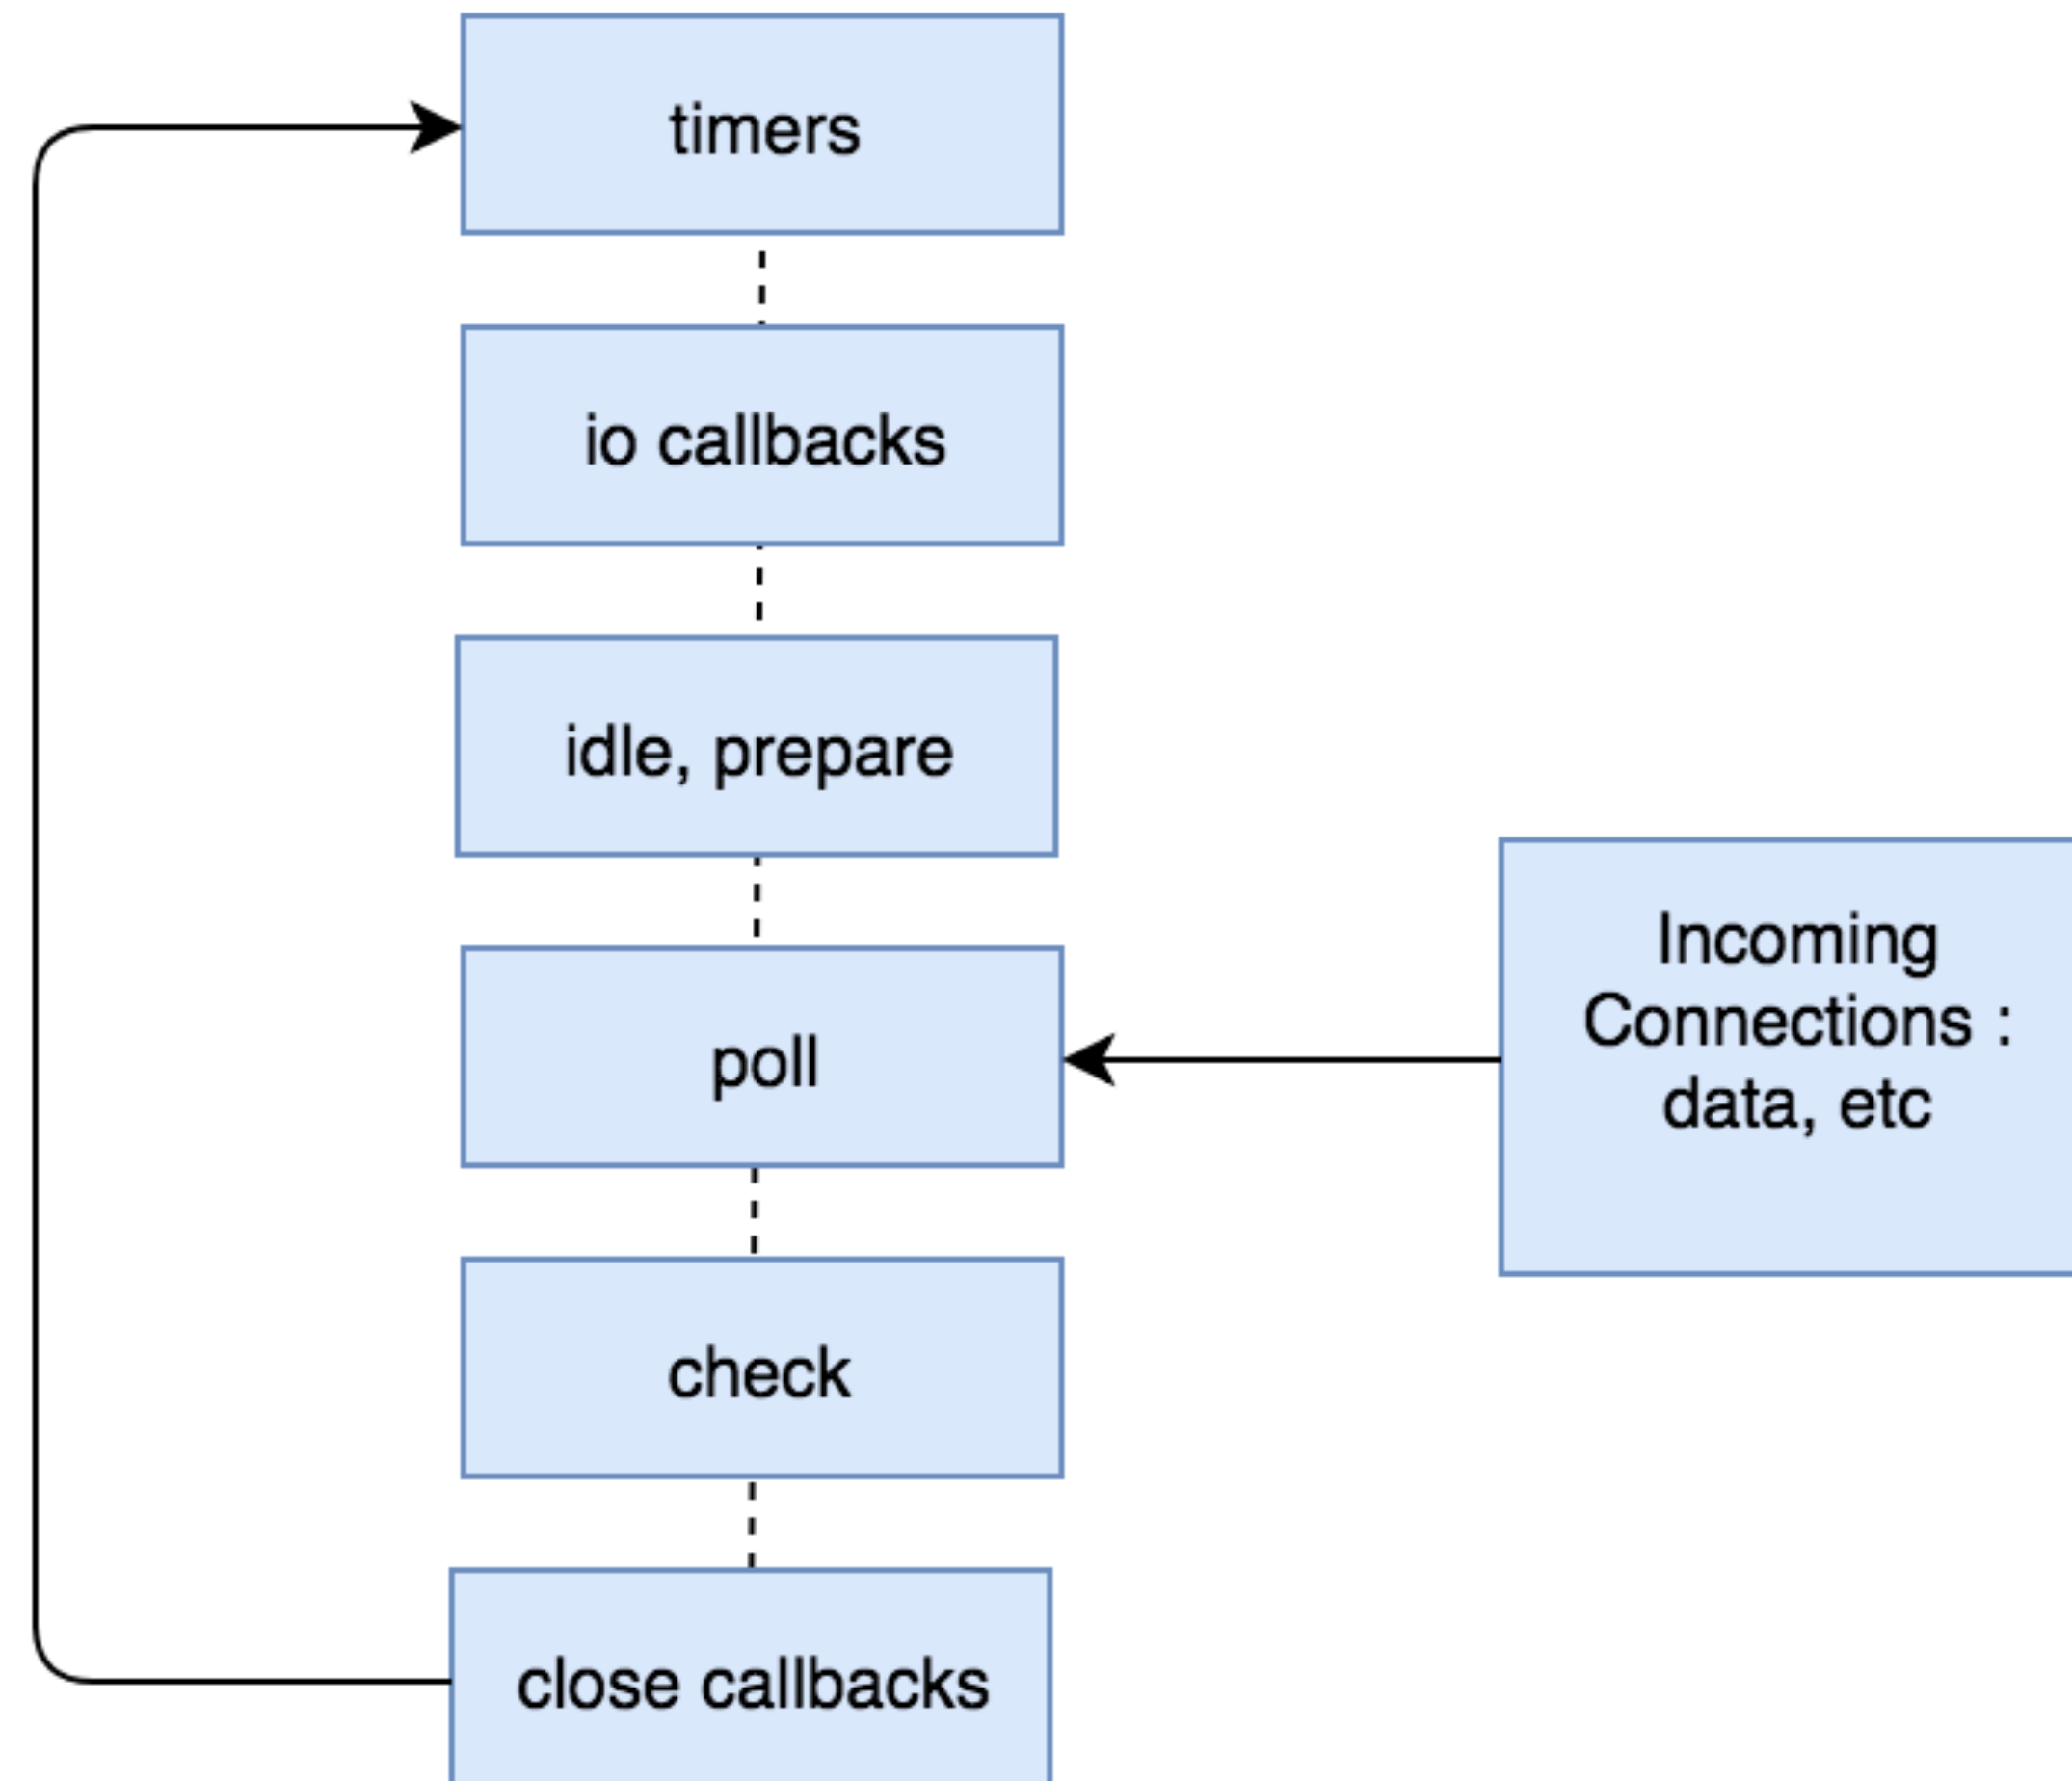


libuv

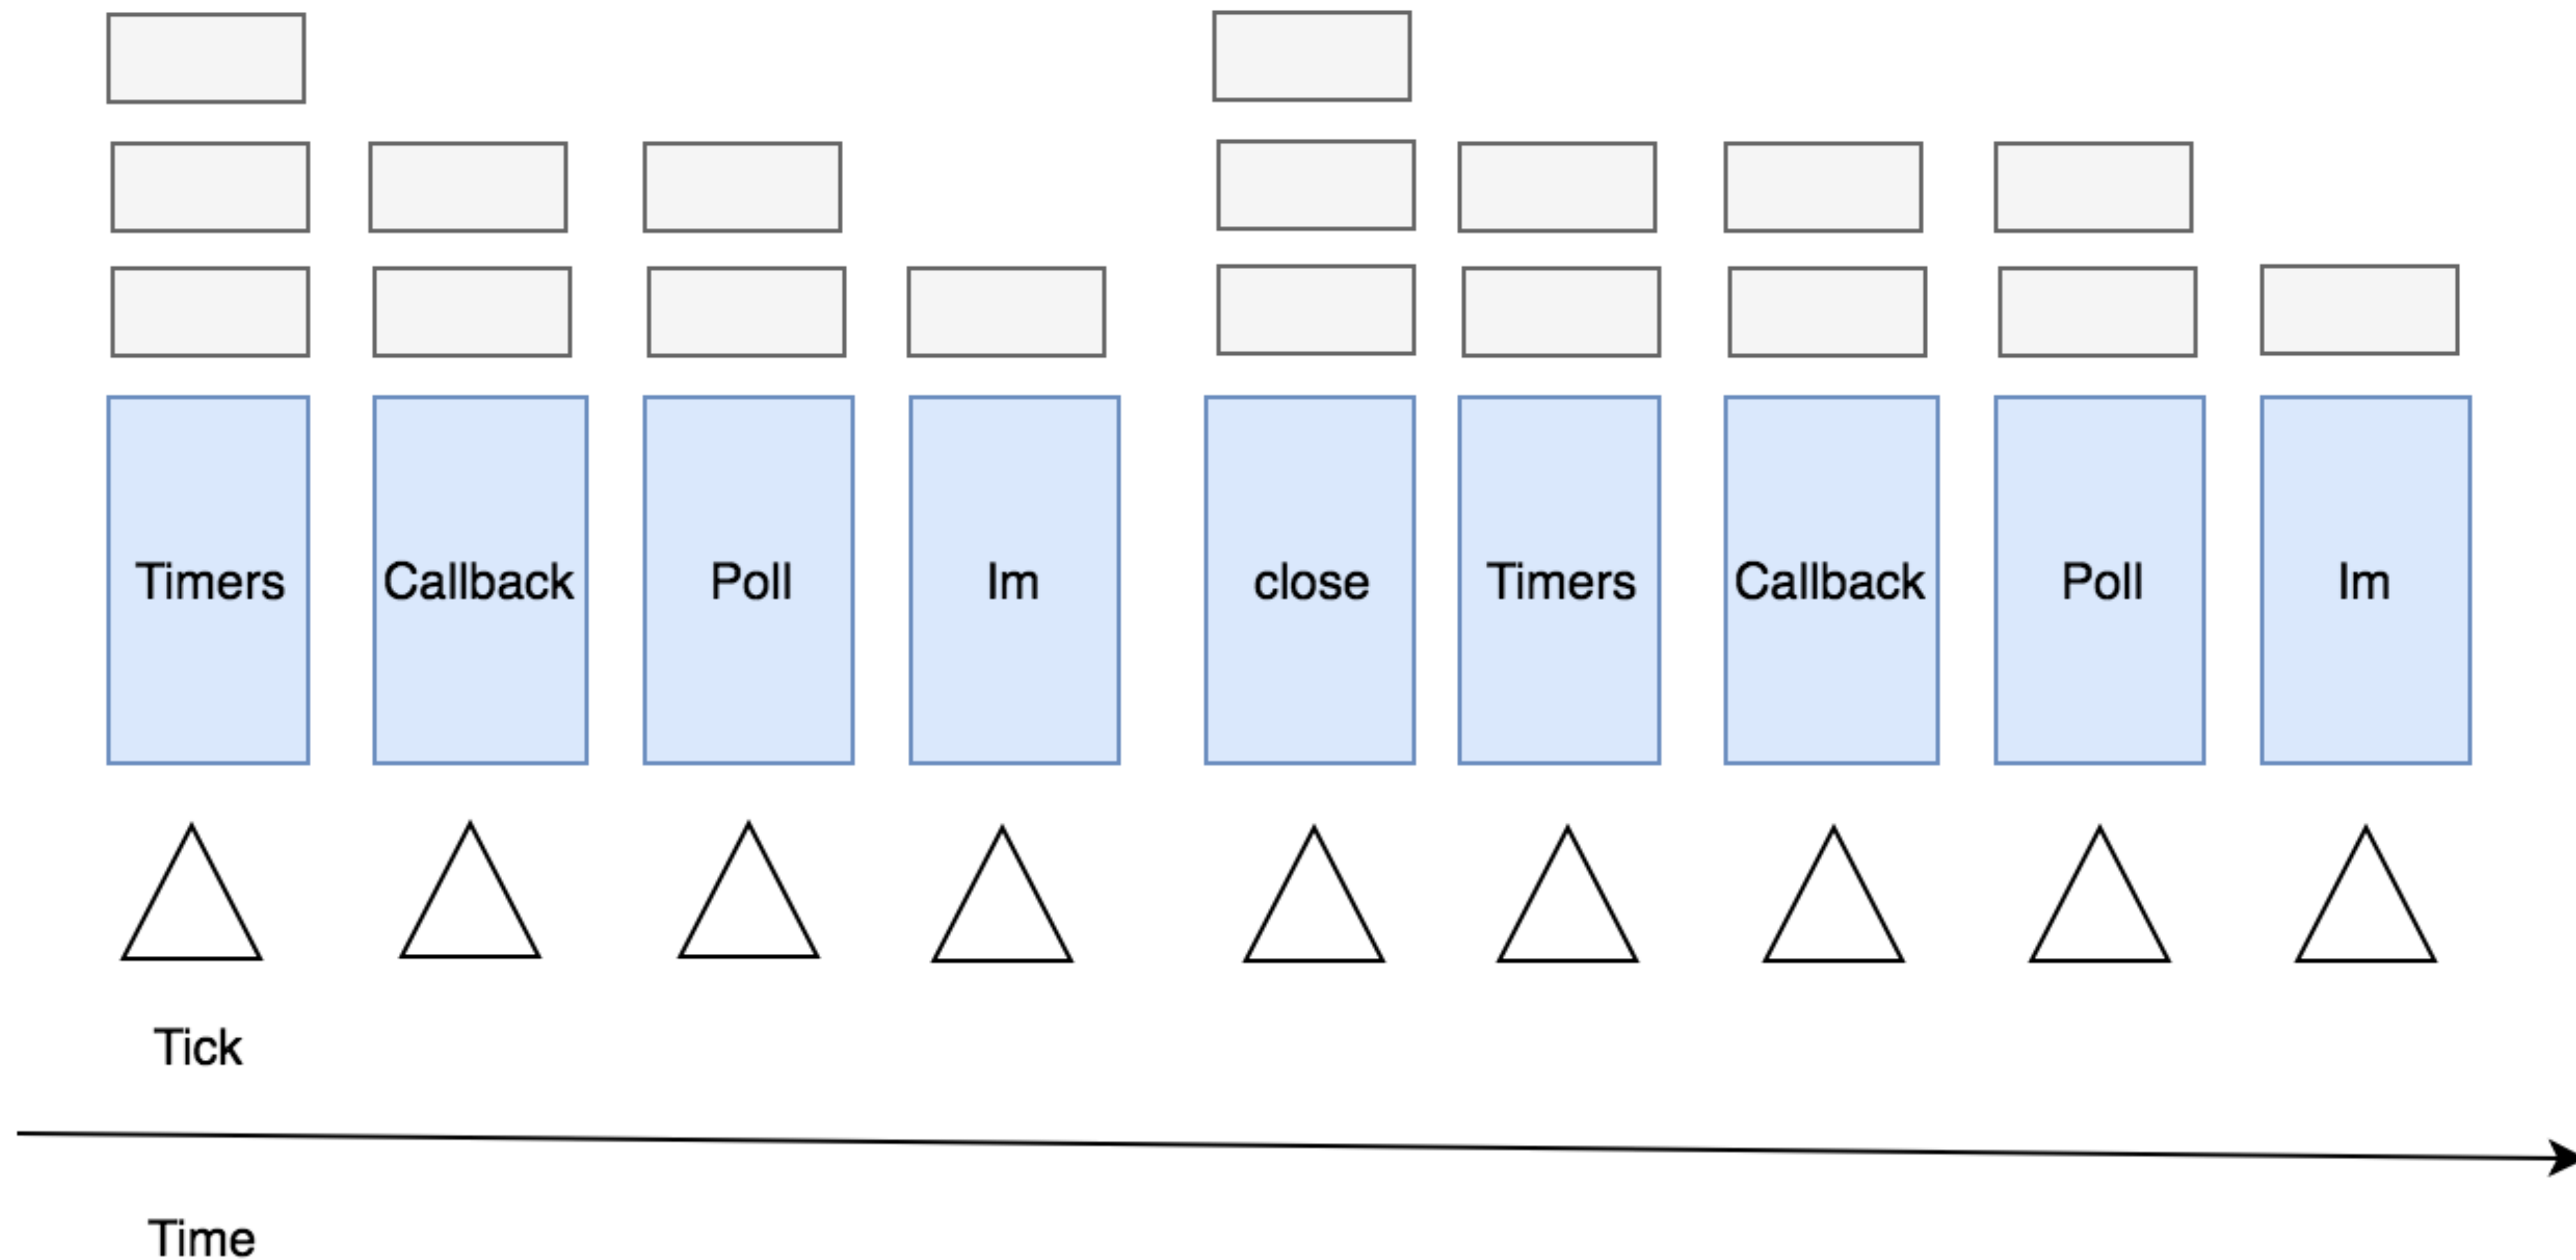
- Libuv in a nutshell :
 - Multi platform C library
 - Provides support for async I/O based on event loop
- Supports :
 - Epoll (Linux)
 - Kqueue (OSX)
 - Windows IOCP
 - Solaris event ports

The Event Loop - The Different Phases

Event Loop Phases Overview



Phase General Mechanism



Timers Phase

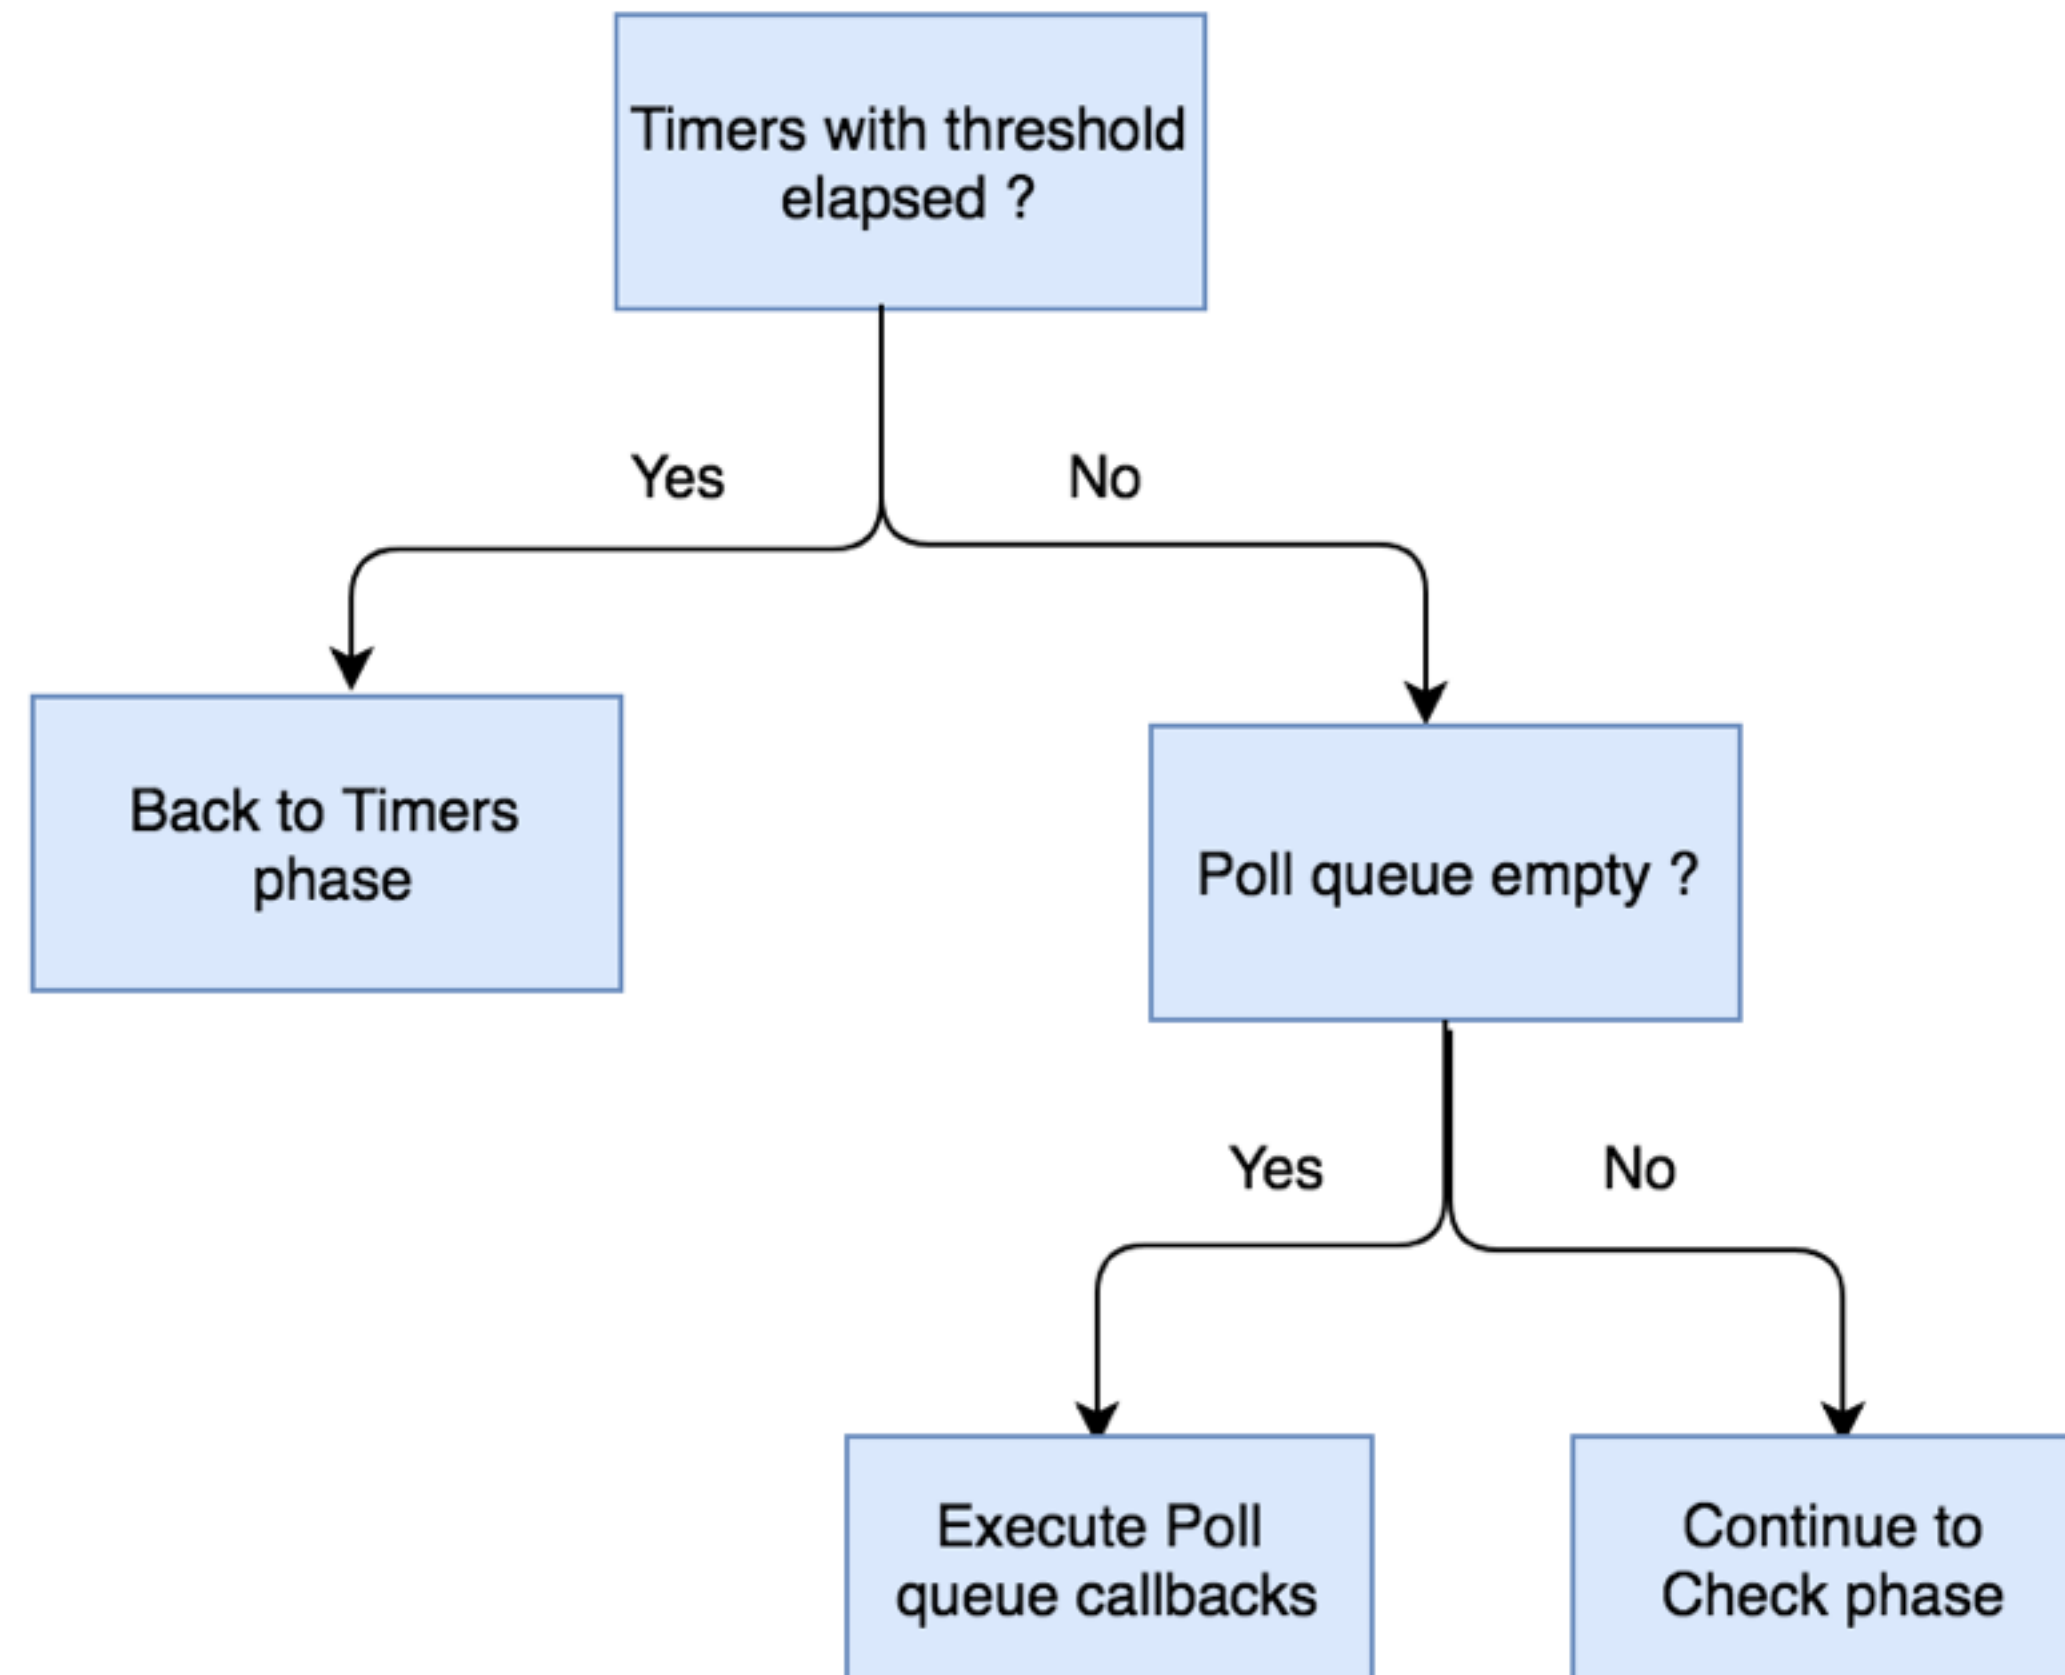
- `setTimeout`, `setInterval`
- Timer's callback will run as soon as they can be scheduled after the threshold
- Timer's callback scheduling controlled by the “poll” phase

I/O Callback Phase

- Executes system error callbacks
 - Example : TCP socket connection error.
- Normal I/O operation callbacks are executed in the poll phase.

```
socket.on('error', function(exception) {  
  console.log('SOCKET ERROR');  
  socket.destroy();  
})
```

Poll phase



Check Phase And Close Phase

- Check Phase - Execute callbacks for setImmediate timers
- Close Phase - Handles an abruptly close of a socket or a handle


Lets profile some code

The JIT Compiler And V8 Engine

What Is Just-In-Time Compilation ?

- Compilation during run time
- Combines two approaches :
 - Ahead of compilation
 - Interpreter


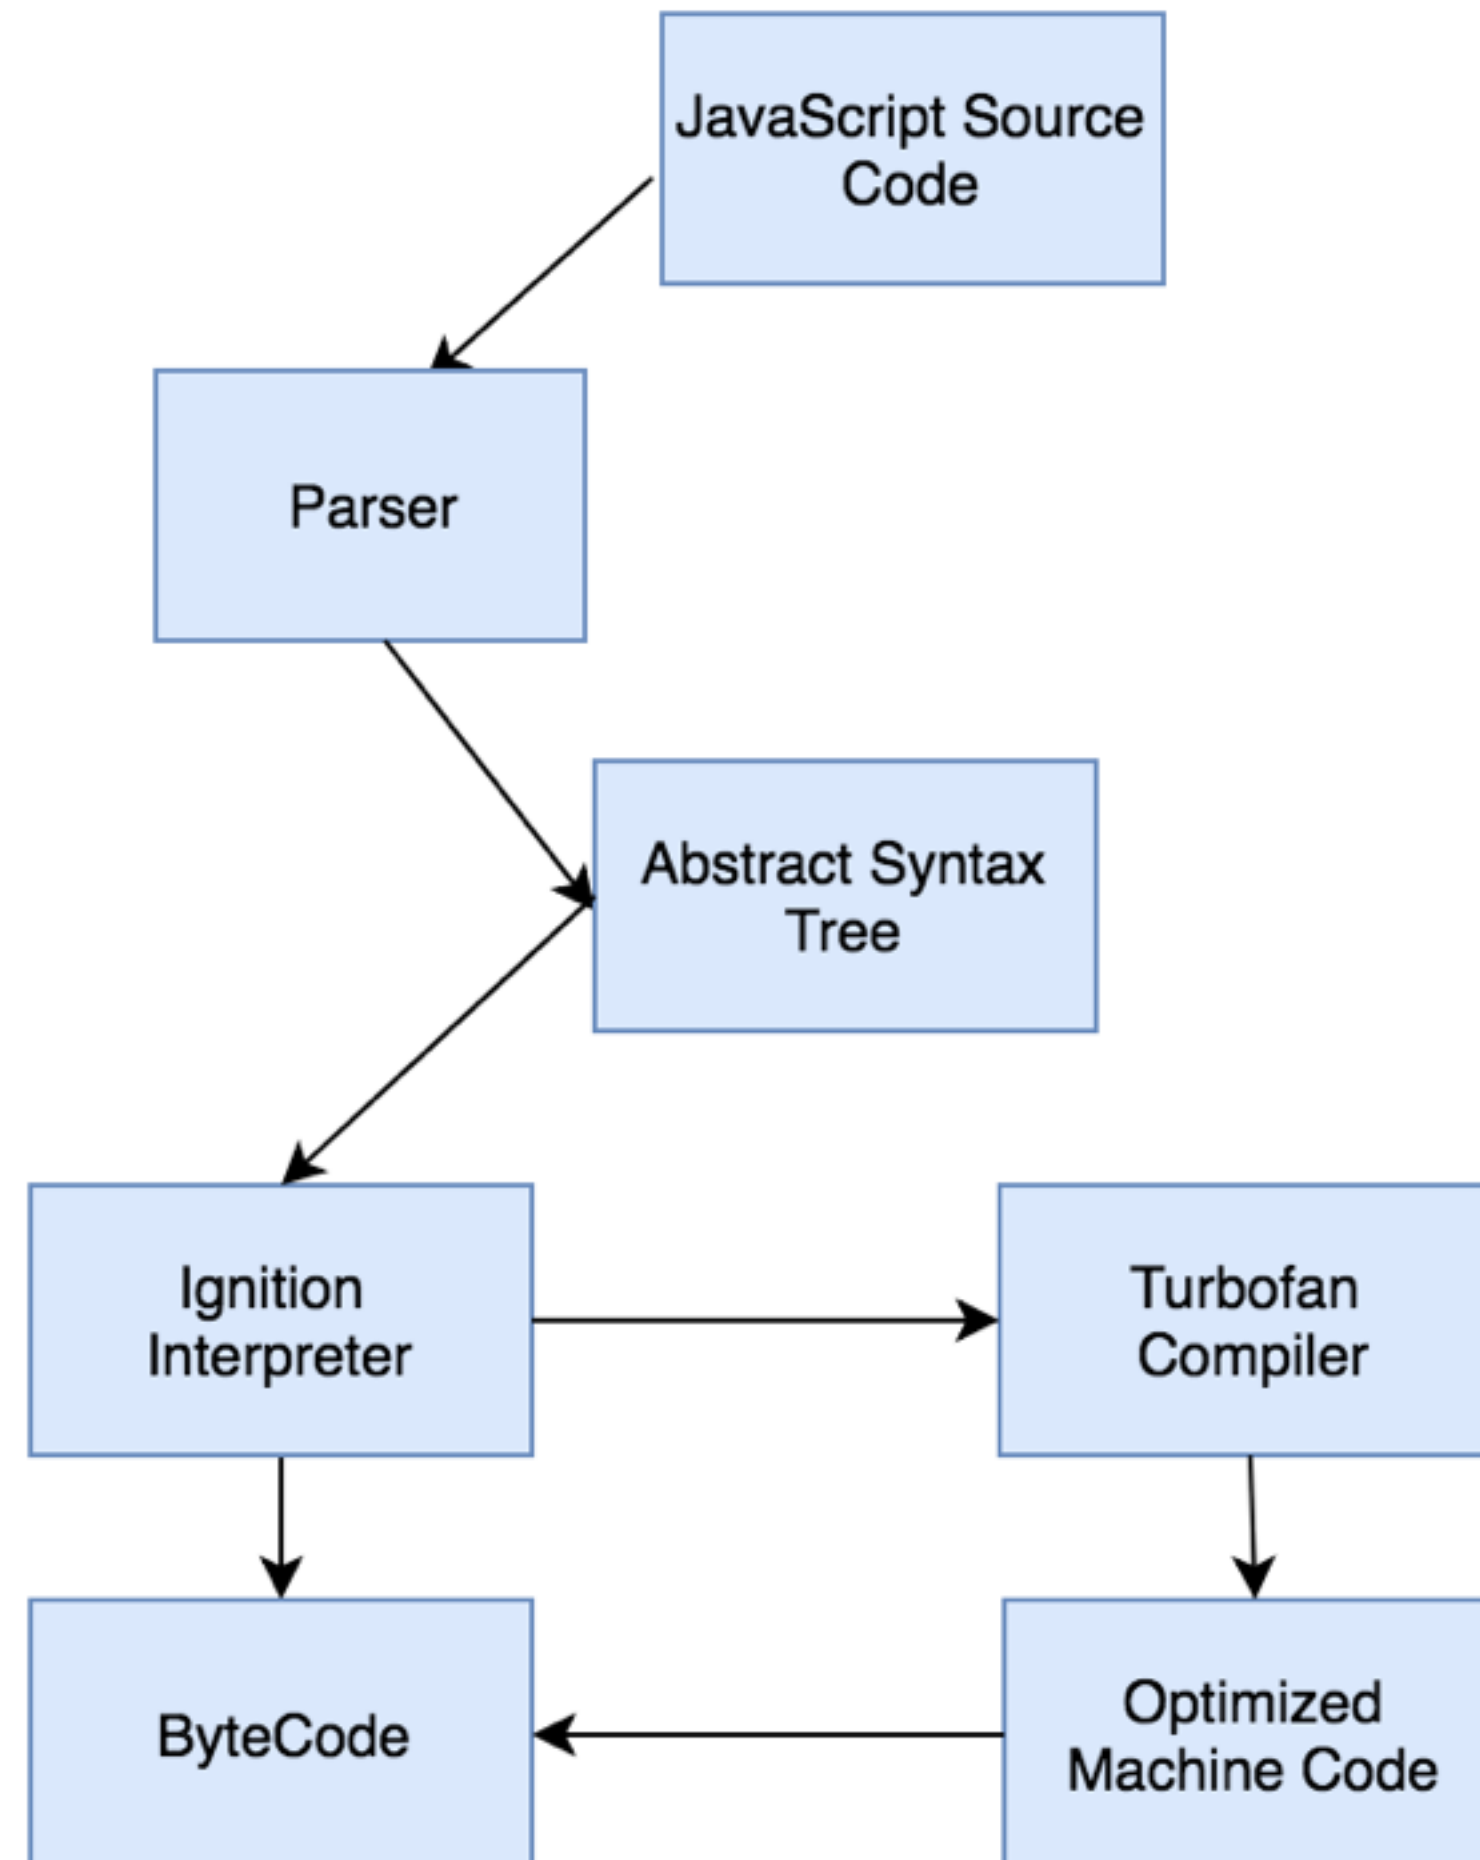
JIT Compiler In Java




Chrome

- Open source JavaScript engine
- Developed originally for Google Chrome and chromium
- Also used for
 - Couchbase
 - MongoDB
 - Node.js

V8 Architecture




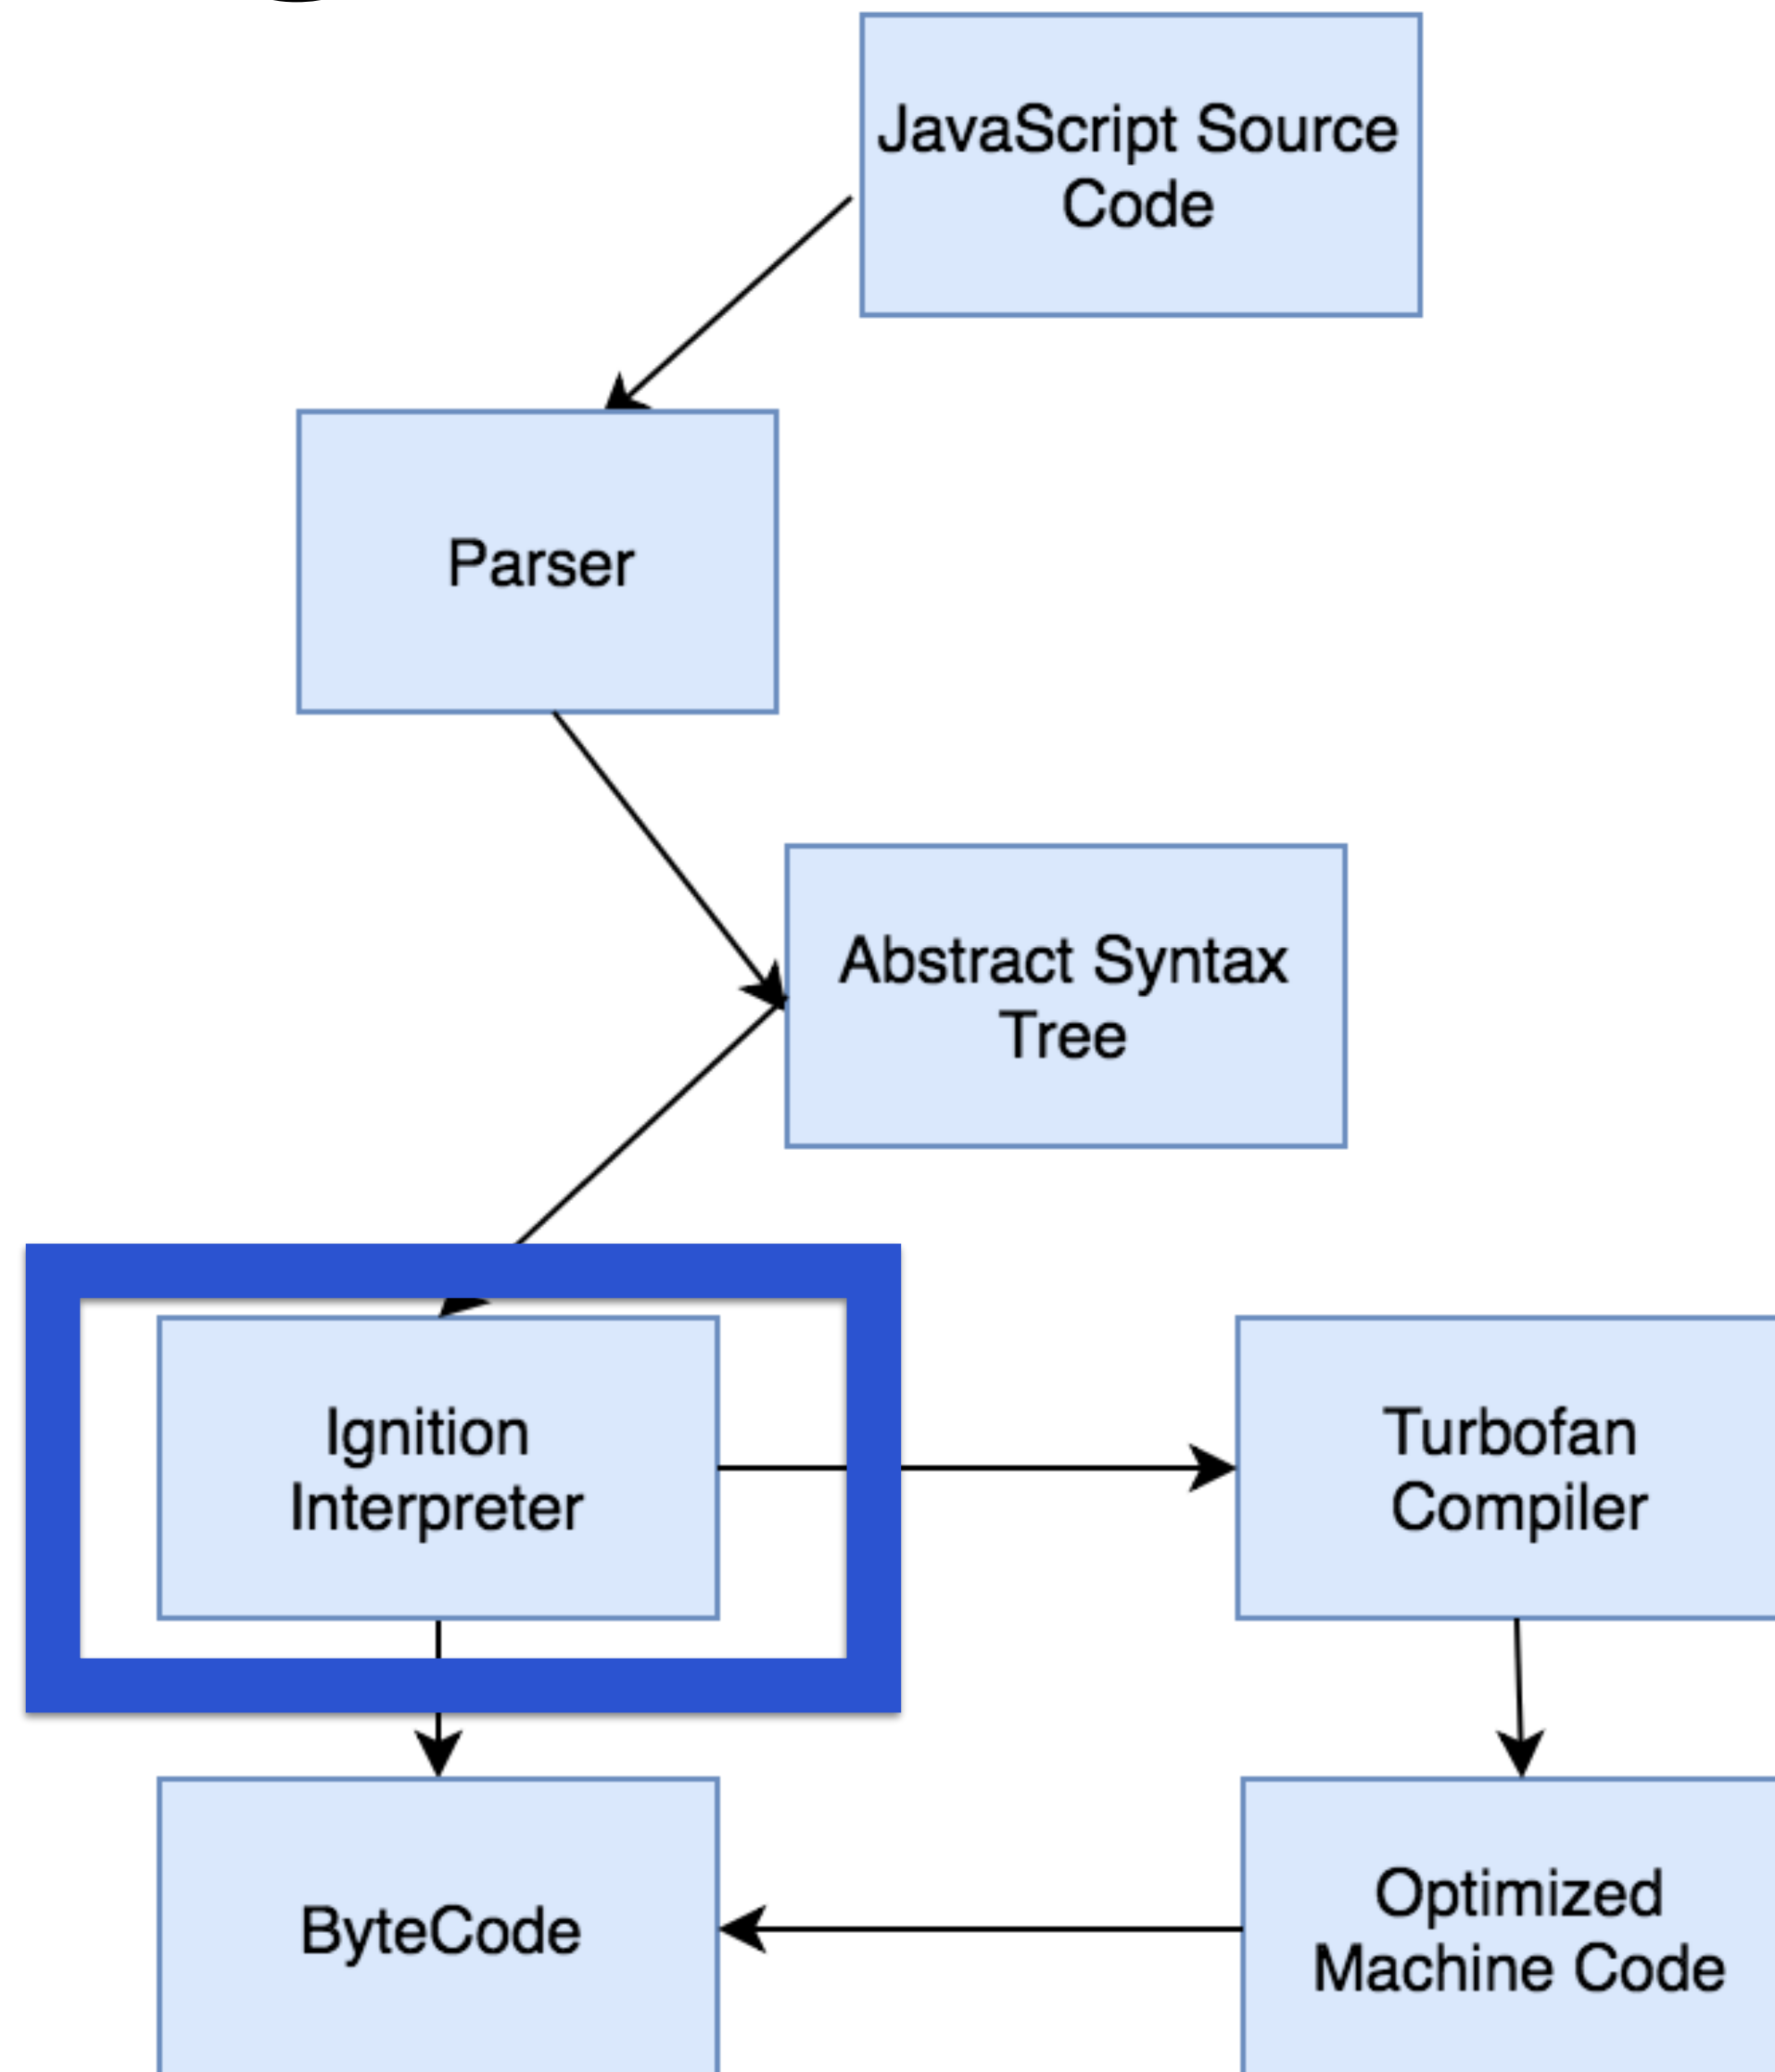
The JIT Compilation In V8




What Is An Optimised Compiler?

- When a code is hot - it is worth doing multiple optimisations
- Tracing will send it to optimising compiler
- Creates an even faster version of the code
- Tracing pulls it when the function code runs

Ignition Interpreter

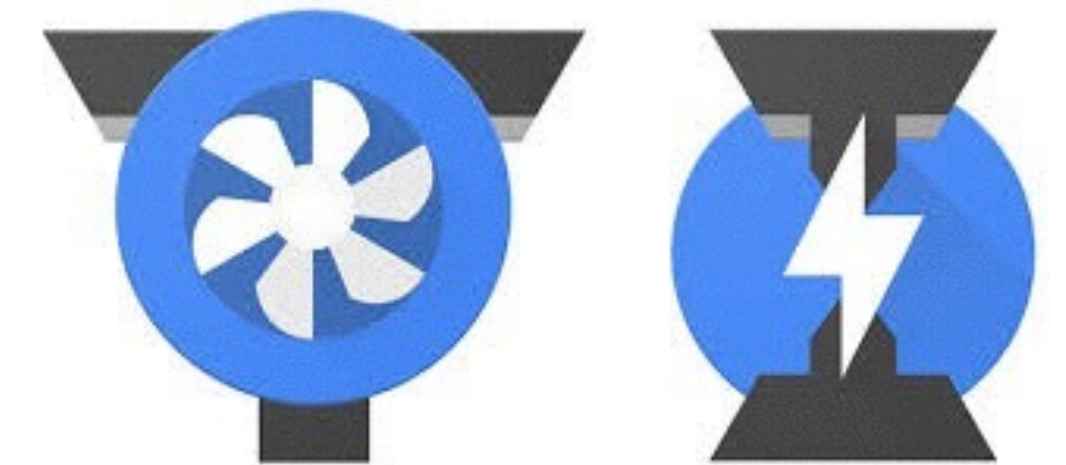



Ignition Interpreter




- Interpreter for V8
- Translates to low level Bytecode
- Enabling the following code to be stored more compactly in Bytecode
 - Run once code
 - Non hot code

Turbofan Compiler



Turbofan Compiler




- Make hot code run as fast as possible
- Relies on input type information collected via inline caches while functions run via the Ignition interpreter.
- Generates the best possible code handling the different types it encountered
- The fewer function input type variations the compiler has to consider, the smaller and faster the resulting code will be

How To Help Turbofan Optimise Hot Code ?

- The fewer function input type variations lead to smaller and faster resulting code.
- Keeping your functions *monomorphic* or at least *polymorphic*
 - *Monomorphic*: one input type
 - *Polymorphic*: two to four input types
 - *Megamorphic*: five or more input types

Optimisation And De-optimisation

- **Optimisation** - All assumptions fulfilled - Compiled code runs.
- **Deoptimisation** - Not all assumptions fulfilled - Compiled code erased



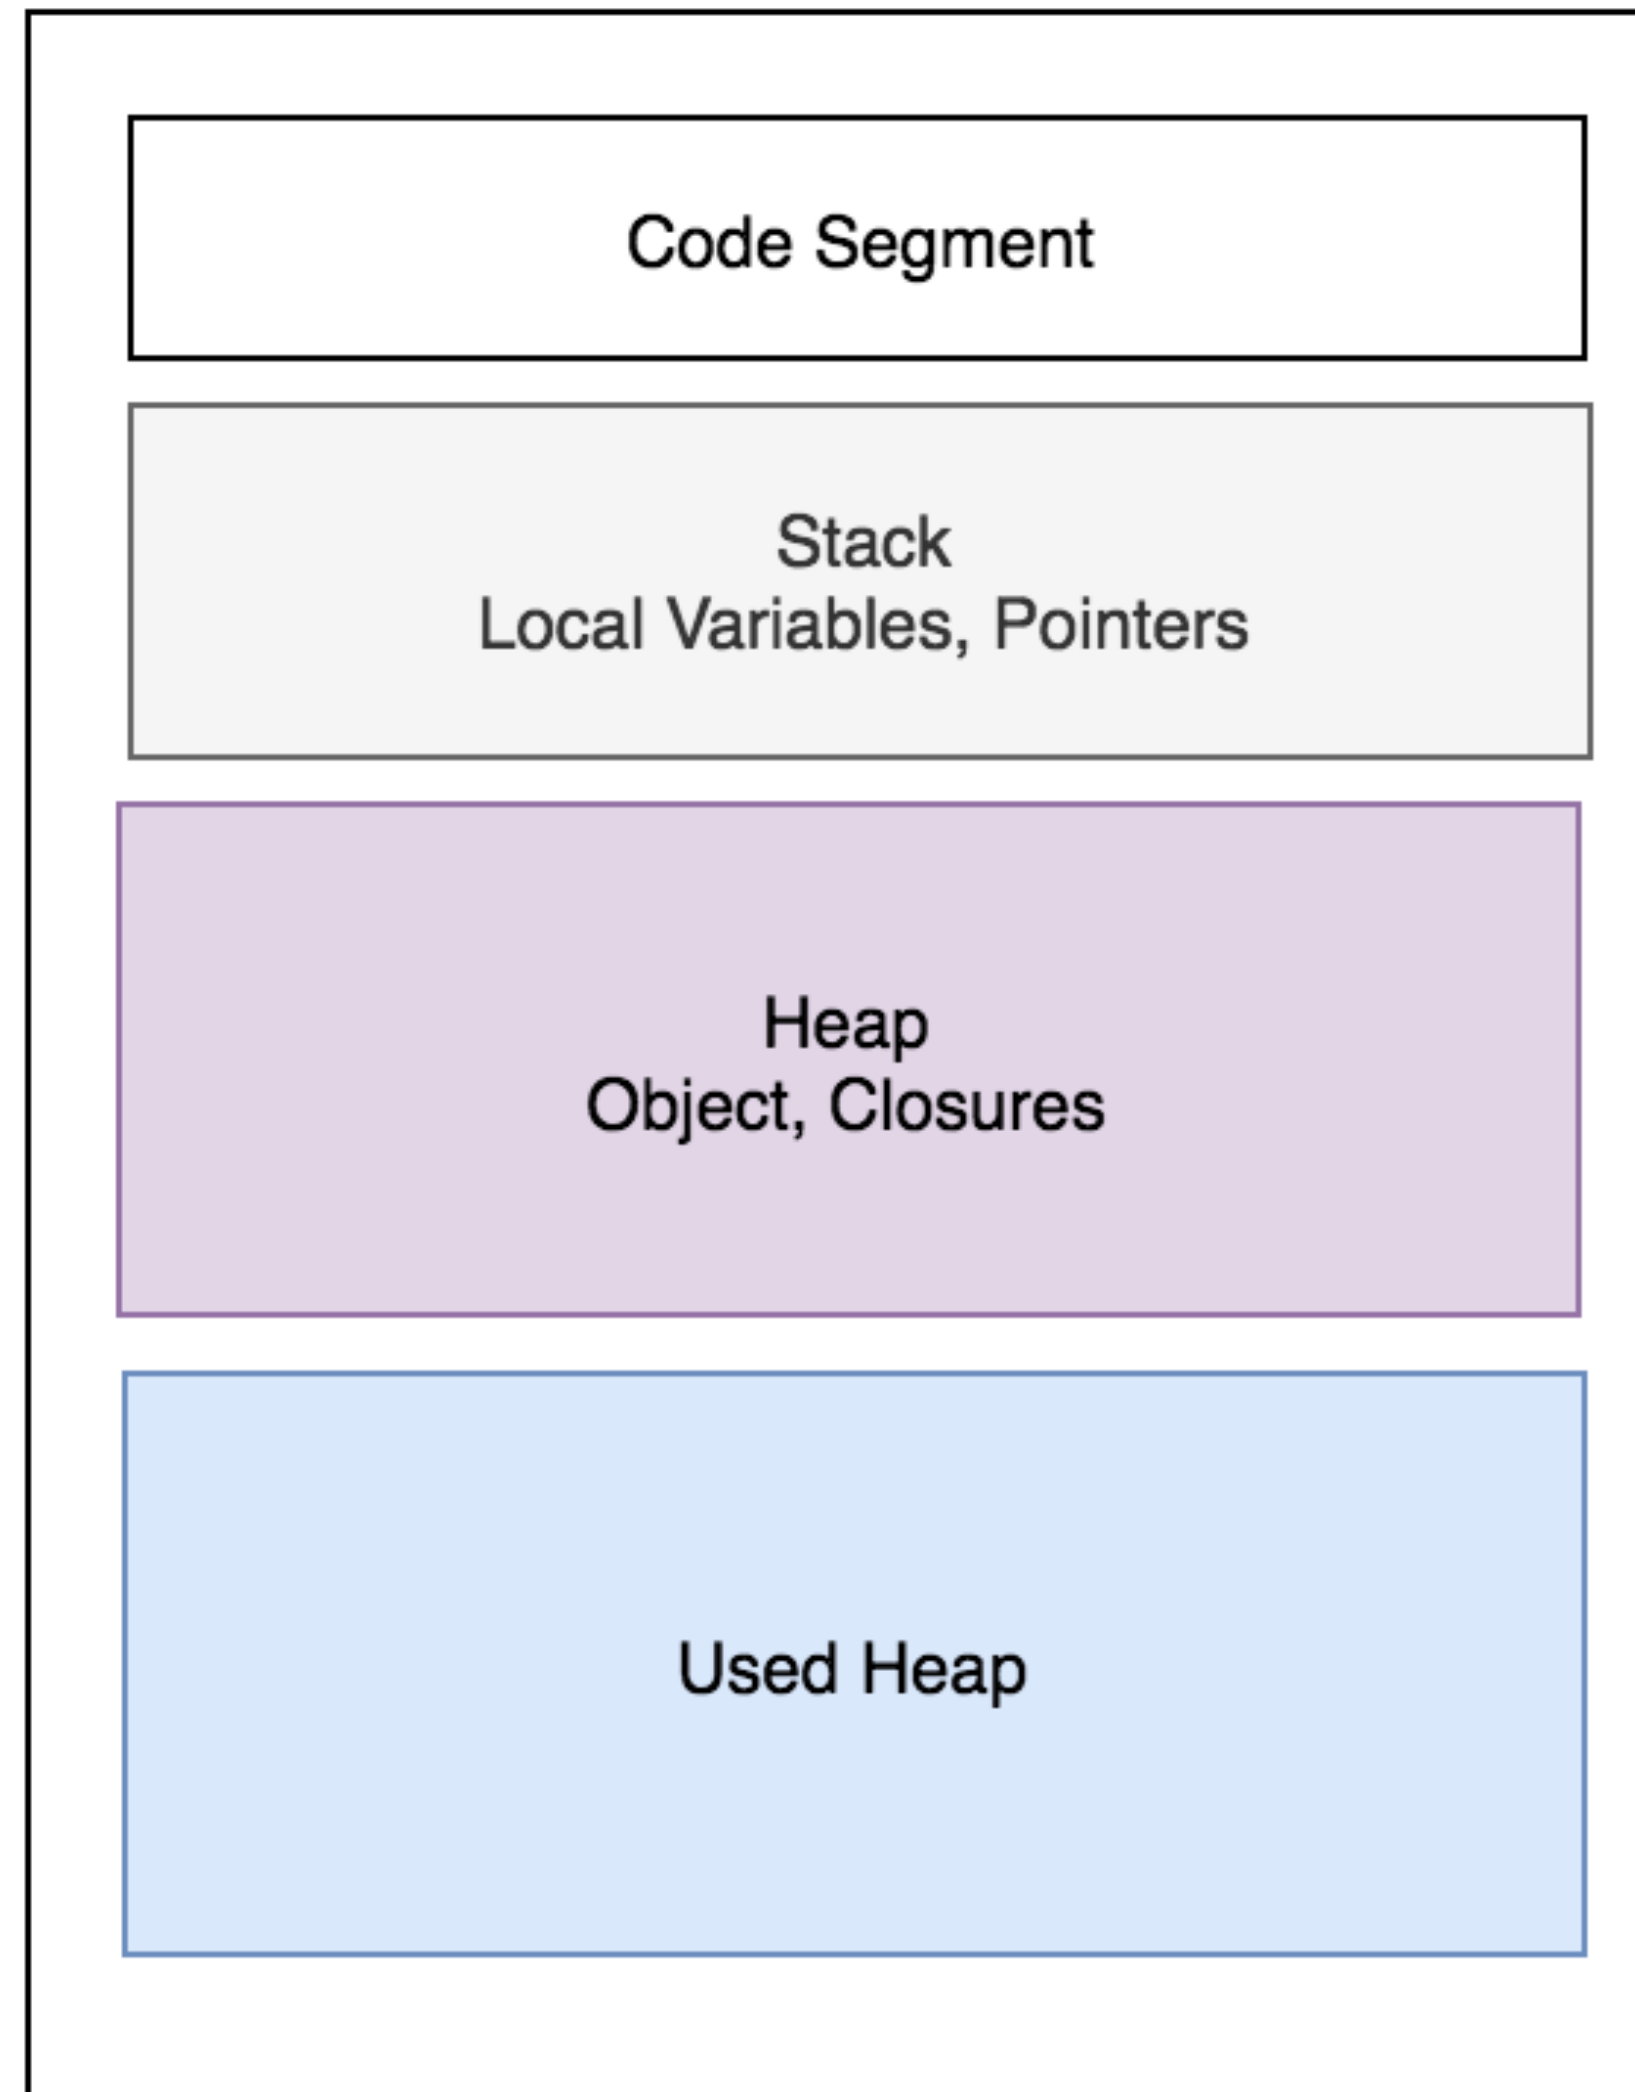
Avoid De-optimisation

- When a code is optimised and de-optimised - it ended up being slower then just use the baseline compiled version
- Most browsers and engines will stop trying after several iterations of optimising and de-optimizing

De-optimisation Demo

V8 Memory Management


V8 Memory Structure



The Stack

- Every executing function pushes its local variables and arguments
- Maintains two pointers - Stack Pointer and Base Pointer
- Divided into stack frames
- No Garbage Collection - self cleaning
- Hold the key to the garbage collection process - Starts from active objects that has pointers to the stack.

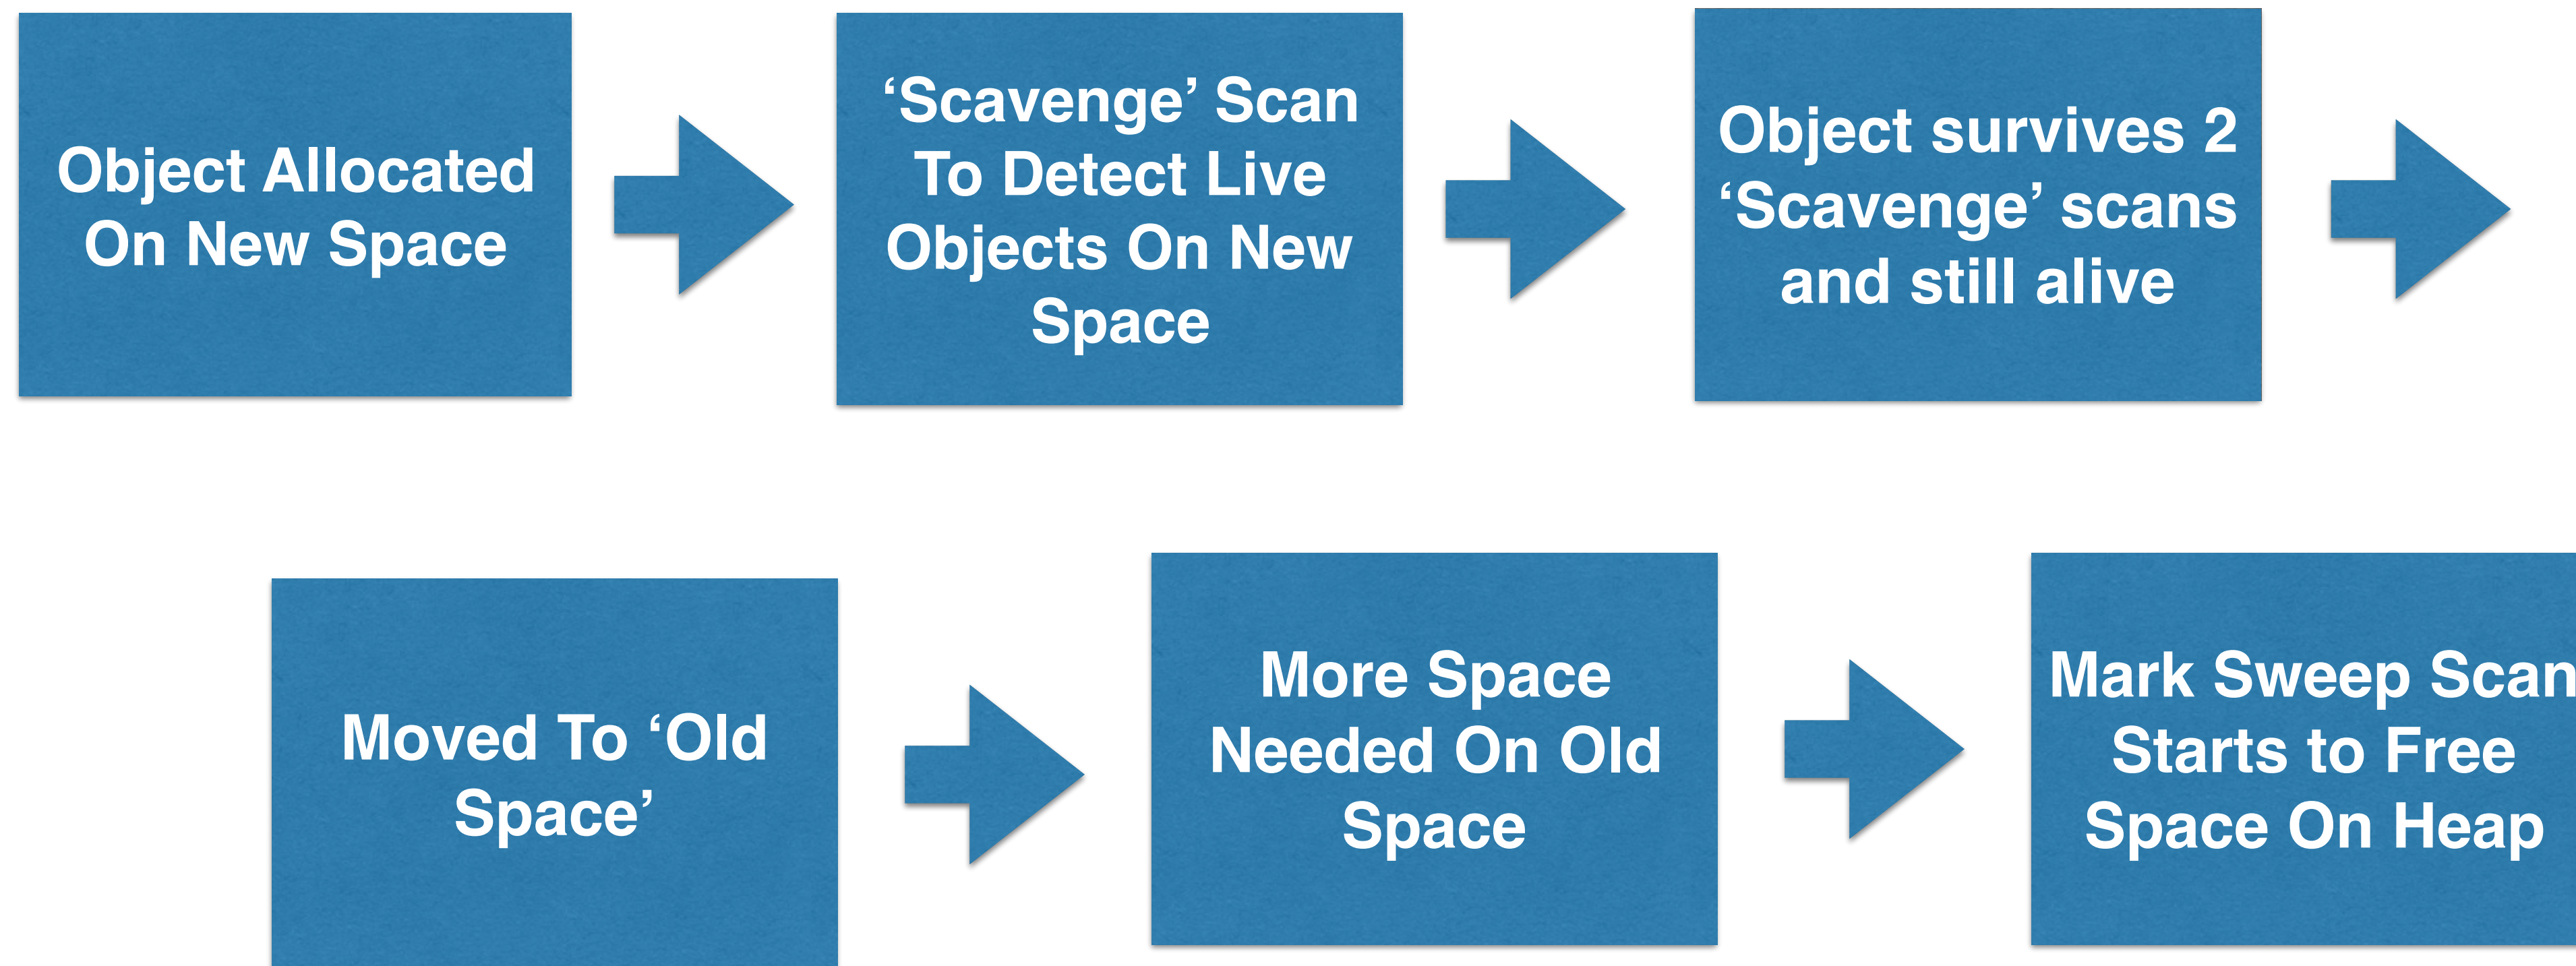
The Heap Structure



Generational GC System



- Lifetime of objects determines their place on the heap
- New Space - short lived Objects
- Old Space - long lived Objects
- Scan on 'New Space' called 'Scavenge'
 - Very fast - takes less than a millisecond
- Scan on 'Old Space' called 'Mark Sweep'
 - Slower scan

Object Life On The Heap





Scavenge Scan Explained


To space filled,
Scavenge triggered



All objects move to from
space




All live objects move
back to 'To Space'




Mark Sweep For The Following Code

```
const a = {};  
a.b = {};  
a.b.c = [1, 2];  
a.b.d = [3, 4];  
a.b.d = null;
```

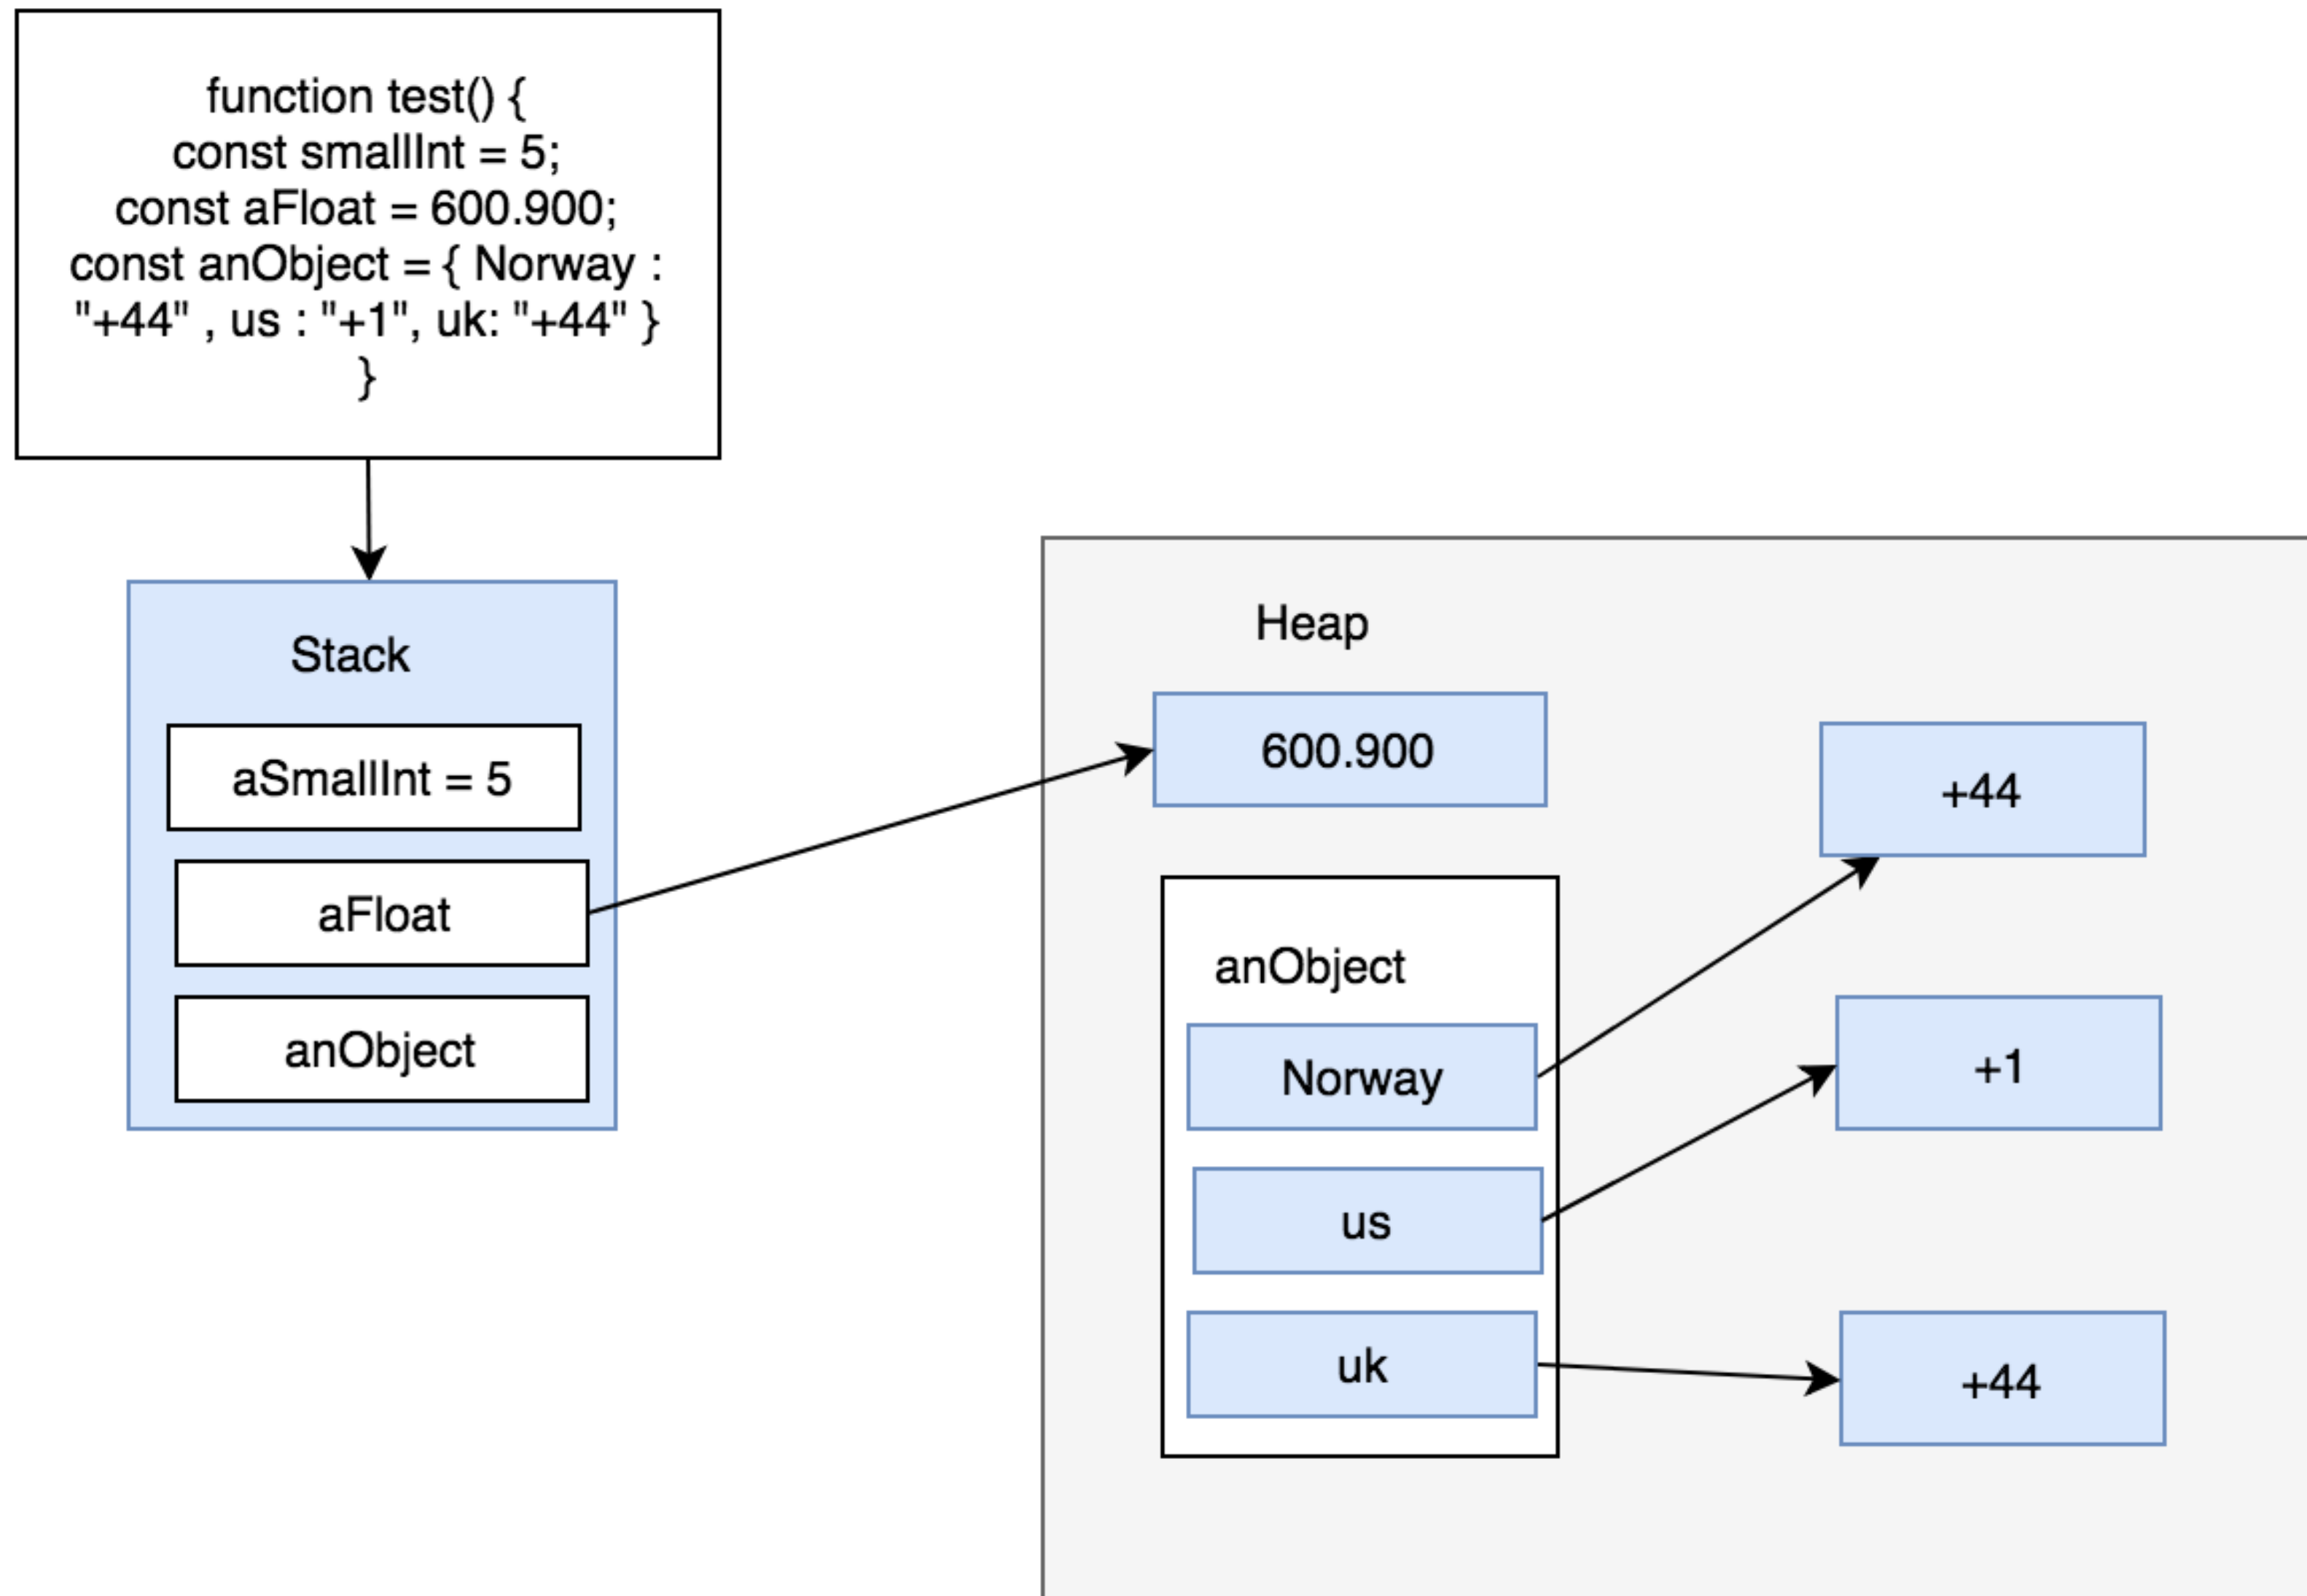

Mark Sweep - DFS



Mark Sweep - DFS



Whats Allocated Where ?



Lets Find Some Memory Leaks !



- Twitter: **@SternTwena**